

Nidra: An Extensible Android Application for Recording, Sharing and Analyzing Breathing Data Collected During Sleep

Jagat Deep Singh



Thesis submitted for the degree of
Master in Programming and Networks
60 credits

Department of Informatics
Faculty of Mathematics and Natural Sciences

UNIVERSITY OF OSLO

Autumn 2019

Nidra: An Extensible Android Application for Recording, Sharing and Analyzing Breathing Data Collected During Sleep

Jagat Deep Singh

© 2019 Jagat Deep Singh

Nidra: An Extensible Android Application for Recording, Sharing and
Analyzing Breathing Data Collected During Sleep

<http://www.duo.uio.no/>

Printed: Reprocentralen, University of Oslo

Acknowledgements

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Leo a diam sollicitudin tempor. Elit ullamcorper dignissim cras tincidunt lobortis feugiat vivamus. Consectetur a erat nam at lectus urna duis. Ullamcorper sit amet risus nullam eget felis eget nunc. Sollicitudin nibh sit amet commodo. Volutpat maecenas volutpat blandit aliquam etiam erat. Sed viverra ipsum nunc aliquet bibendum enim facilisis gravida neque. Metus dictum at tempor commodo. Nunc vel risus commodo viverra maecenas accumsan lacus. Vitae justo eget magna fermentum iaculis eu non diam. Habitant morbi tristique senectus et. Sed enim ut sem viverra aliquet. Lectus mauris ultrices eros in cursus.

Turpis massa tincidunt dui ut ornare lectus. Elit sed vulputate mi sit amet mauris commodo quis imperdiet. Etiam non quam lacus suspendisse faucibus interdum posuere lorem ipsum. Morbi non arcu risus quis. Quis viverra nibh cras pulvinar mattis nunc sed. Tellus cras adipiscing enim eu turpis egestas. Nec tincidunt praesent semper feugiat nibh sed. Ipsum dolor sit amet consectetur. Duis convallis convallis tellus id interdum. Nulla aliquet enim tortor at.

Abstract

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Leo a diam sollicitudin tempor. Elit ullamcorper dignissim cras tincidunt lobortis feugiat vivamus. Consectetur a erat nam at lectus urna duis. Ullamcorper sit amet risus nullam eget felis eget nunc. Sollicitudin nibh sit amet commodo. Volutpat maecenas volutpat blandit aliquam etiam erat. Sed viverra ipsum nunc aliquet bibendum enim facilisis gravida neque. Metus dictum at tempor commodo. Nunc vel risus commodo viverra maecenas accumsan lacus. Vitae justo eget magna fermentum iaculis eu non diam. Habitante morbi tristique senectus et. Sed enim ut sem viverra aliquet. Lectus mauris ultrices eros in cursus.

Turpis massa tincidunt dui ut ornare lectus. Elit sed vulputate mi sit amet mauris commodo quis imperdiet. Etiam non quam lacus suspendisse faucibus interdum posuere lorem ipsum. Morbi non arcu risus quis. Quis viverra nibh cras pulvinar mattis nunc sed. Tellus cras adipiscing enim eu turpis egestas. Nec tincidunt praesent semper feugiat nibh sed. Ipsum dolor sit amet consectetur. Duis convallis convallis tellus id interdum. Nulla aliquet enim tortor at..

Contents

List of Listings	vi
List of Tables	vii
List of Figures	ix
I Introduction and Background	1
1 Introduction	3
1.1 Background and Motivation	3
1.2 Problem Statement	3
1.3 Limitations	5
1.4 Research Methods	5
1.5 Contributions	5
1.6 Thesis Outline	5
2 Background	7
2.1 CESAR: Project Structure	7
2.1.1 Extensible Data Acquisition Tool	7
2.1.2 Extensible Data Stream Dispatching Tool	9
2.1.3 Flow Sensor Kit	12
2.2 Android OS	13
2.2.1 Android Architecture	13
2.2.2 Application Components	14
2.2.3 Process and Threads	18
2.2.4 Inter-Process Communication (IPC)	19
2.2.5 Data and File Storage	19
2.2.6 Architecture Patterns	20
2.2.7 Power Management	21
2.2.8 Bluetooth Low Energy	22
3 Related Work	23
II Design and Implementation	25
4 Analysis and High-Level Design	27

4.1	Requirement Analysis	28
4.1.1	Stakeholders	28
4.1.2	Resource Efficiency	28
4.1.3	Security and Privacy	28
4.2	High-Level Design	29
4.2.1	Task Analysis	29
4.3	Separation of Concerns	31
4.3.1	Recording	31
4.3.2	Sharing	34
4.3.3	Modules	36
4.3.4	Analytics	37
4.3.5	Storage	38
4.3.6	Presentation	40
4.4	Data Structure	42
4.4.1	Data Formats	42
4.4.2	Data Entities	44
4.4.3	Data Packets	47
5	Implementation	51
5.1	Android Components	51
5.1.1	Inter-Process Communication: AIDL	51
5.2	Flow Sensor Kit Development	53
5.3	Implementation of Concerns	53
5.3.1	Recording	53
5.3.2	Sharing	58
5.3.3	Modules	61
5.3.4	Analytics	63
5.3.5	Storage	65
5.3.6	Presentation	66
III	Evaluation and Conclusion	71
6	Evaluation	73
6.1	Experiments and Measurements	73
6.1.1	Experiment A: Orchestral Concert to Analyze Musical Absorption using Nidra to Collect Breathing Data	73
6.1.2	Experiment B: 8-hours recording	76
6.1.3	Experiment C: User-Experience	76
6.1.4	Experiment D: Creating a Simple Module	76
6.2	Main Findings	76
7	Conclusion	77
Appendix		78
A	Power Data	81

Listings

4.1	My Caption	43
4.2	My Caption	43
4.3	My Caption	47
4.4	My Caption	48
5.1	My Caption	53
5.2	My Caption	54
5.3	My Caption	59
5.4	My Caption	60
5.5	My Caption	62

List of Tables

4.1	Example entry in record table	45
4.2	Example entry in sample table	46
4.3	Example entry in module table	46
6.1	Device models used during the concert	74
6.2	Day 1—Duration: 1 hour & Roof Samples: 5145	76
6.3	Day 2—Duration: 50 mins. & Roof Samples: 4288	76

List of Figures

1.1	Structure of the project, separating functionality into three independent layers [9]	4
2.1	Sharing the collected data between multiple applications [22]	10
2.2	Sharing the collected data between multiple applications [9]	11
2.3	Recording	13
2.4	Recording	16
2.5	Entity Relationship Diagram	20
4.1	Recording	29
4.2	Recording	32
4.3	Sharing	35
4.4	Modules	36
4.5	Analytics	38
4.6	Storage	39
4.7	Presentation	41
4.8	Modules	44
5.1	Applications components	52
5.2	Implementation of recording functionality: (A) Start a Recording	54
5.3	Implementation of recording functionality: (B) Stop a Recording	56
5.4	Implementation of recording functionality (C)	57
5.5	Implementation of sharing functionality (A): Exporting one or all Records	58
5.6	Implementation of sharing functionality (B)	60
5.7	Implementation of module functionality(A): Add a Module	62
5.8	Implementation of module functionality(B): Launch a Module	63
5.9	Implementation of analytics functionality (A): Display a Graph for a Single Record	64
5.10	Entity Relationship Diagram	65
5.11	The recording screen displayed to the user; with the screen: (A) during a recording, (B) real-time analytics, and (C) finalizing the recording	67
5.12	The sharing screen displayed to the user	68
5.13	The module screen displayed to the user	69
5.14	The analytics screen displayed to the user	69

6.1	Concert Day 1—Mobile Device B	75
6.2	Concert Day 2—Mobile Device B	75

Part I

Introduction and Background

Chapter 1

Introduction

1.1 Background and Motivation

The CESAR project aim to use low cost sensor kit to prototype applications using physiological signals related to heart rate, brain activity, oxygen level in blood to monitor sleep and breathing related illnesses, like Obstructive Sleep Apnea (OSA). Side effects of OSA do not only cause sleepiness during day time (which might affect daily chores), but also serious illnesses like diabetes and cardiac dysfunctions. Statistically speaking, it is estimated that about 25% of the adult population in Norway has OSA, but only 10% of them are diagnosed. A major problem with diagnosing OSA is polysomnography in *sleep laboratories* [41]. This is both really expensive and inefficient due to lacking capacity to perform sufficient tests with patients. Hence, the CESAR project aims to contribute to this situation with a low-cost Android and BiTalino based system to tackle these problems in a minimal invasive approach.

The project has been developed by various people over the years, and the system has been divided into three parts (illustrated in Figure 1.1). The data acquisition part, the data streams dispatching part, and the application part. The first two parts are already implemented (summarized in the section below), thus, the last part is what we will be focusing throughout this thesis.

1.2 Problem Statement

As indicated in the background and motivation section, we decided to look into there are opportunity of extending the system even further. The market has new and affordable sensors that can aid with the data acquisition, which we seamlessly can integrate with the extensible data acquisition tool. In the end, we can hopefully strengthen the detection of

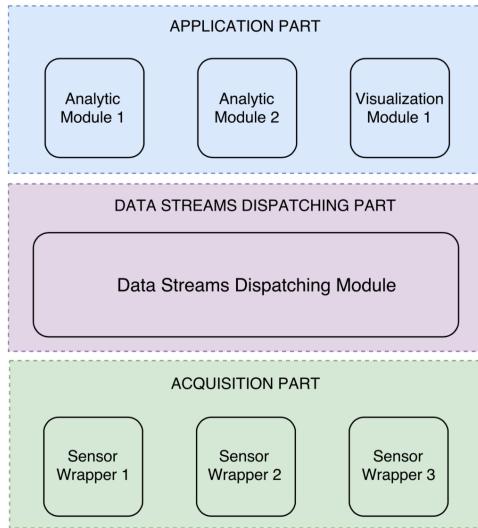


Figure 1.1: Structure of the project, separating functionality into three independent layers [9]

abnormal sleeping patterns, and decreasing the risk of the symptoms they may endure.

In this thesis, we will continue to build on the project by implementing the following:

1. *Integrate new a sensor wrapper*

The following section Flow Sensor Kit introduces to a new sensor type to be used in the CESAR project. Thus, creating a new wrapper to integrate it in the system is necessary. By utilizing the library that comes with the sensor, we can integrate that with the interface and protocol of the CESAR project.

2. *Visualization of the activities on an Android device*

Designing, architecting, and modulating an Android application that implements the new sensor data in an adequately layout.

3. *Detecting anomalies with the help of Machine Learning*

Classifying the data so we can detect abnormal sleeping patterns, and hopefully training a model that can detect the sleeping patterns, locally on the device, without any external supervisor (human intervention) analyzing the data.

1.3 Limitations

1.4 Research Methods

1.5 Contributions

1.6 Thesis Outline

Chapter 2

Background

2.1 CESAR: Project Structure

The CESAR project is a collaboration between the University of Oslo and Oslo Universitetssykehus, with a goal to reduce the threshold to perform a clinical diagnosis of Obstructive Sleep Apnea (OSA) and to reduce the time to diagnose the disorder. Obstructive Sleep Apnea is a common sleeping disorder which affects the natural breathing cycle by reducing respiration or all airflow. As a consequence, OSA can lead to serious health implications, and in some cases, death through suffocation [10].

The CESAR project aims to use low cost sensor kit to prototype applications using physiological signals related to heart rate, brain activity, oxygen level in blood to monitor sleep and breathing related illnesses (e.g., Obstructive Sleep Apnea). As of now, the development in the project facilities for data acquisition from various sensor sources and delegation of the data to subscribed applications. The Section X and Section X is a summary of the development and results of the development, and Figure X is an illustration of the projects development pipeline.

Moreover, the previous work has incorporated support for a few sensor sources (i.e., Bitalino) in order to collect physiological signals. With the support for more sensor sources enables more precise detection and analysis of the disorder can be made. In Section X, a new sensor source is introduced.

2.1.1 Extensible Data Acquisition Tool

In the thesis "Extensible data acquisition tool for Android" by Svein Petter Gjøby [22], we proposed a *data acquisition system* for Android to make application development comprehensible. The thesis proposes a system that hides the low-level sensor specific details into two components, *pro-*

viders and *sensor wrappers*. The provider is responsible for the functionality that is common for all data sources (e.g., starting and stopping the data acquisition), whilst the sensor wrapper is responsible for the data source specific functionality (e.g., communicating with the data source).

The thesis solves the difficulties around creating an extensible data acquisition tool, connecting new and existing sensors, and finding a common interface. The problem statement of the thesis address the following concerns regarding sensors:

- *Common abstraction/interface for the interchanged data*

Sensor platform manufacturers have their own low-level protocol to support the functionality of their product. Typically, the manufacturers provide an software development kits (SDKs) to hide the low-level protocols so third-party development can be easier, however, both the low-level protocol and the SDKs are not standardized. Thus, for each sensor there might be exposure of different commands and methods.

- *Various Link Layer technologies*

Each sensor might use different Link Layer technology (i.e., Ethernet, USB, Bluetooth, WiFi, ANT+ and ZigBee), which means establishing a connection between a device and a sensor might differ. For instance, Bluetooth devices need to be paired, whilst devices on the WiFi can address each other without any pairing.

- *Reusability of sensor code*

Applications that implement support for the low-level protocol of a sensor type can not be shared between different applications. Thus, introducing duplicate work and code if multiple application wish to use the same sensor type. A framework that isolates the sensor that applications can use, might make it easier for application to utilize the collected data. In addition, isolating the sensors into modules improves the robustness and quality of the implementation.

In the thesis, the goal is to develop an extensible system, which enables applications to collect data from various external and built-in sensors through one common interface. The solution around an extensible system is to have the core of the application unchangeable when adding support for a new data source, regardless of the Link Layer technology and communication protocol used by the data source. Making all the data sources behave as the same, is a naive solution to the problem. However, separating the software into two different components, a *provider* component and a *sensor wrapper* component, enables the reuse of functionality that is common amongst the data sources.

The sensor wrapper application is tailored to suit the Link Layer technology and data exchange protocol of one particular data source. Additionally, responsible for connectivity and communication with the data source. The provider application is responsible for managing the sensor wrappers - starting and stopping the data acquisition - and

processing the data received from the sensor wrapper application. Thus, everything that is independent of the data source, should be a part of the provider application. With this type of solution, we gain the possibility to reuse the sensor wrapper application for different provider applications. However, there are some overheads with this solution. Mostly, the interprocess communication that might be costly and increase the complexity of the code. Nonetheless, the flexibility and extensibility gained by the separating the functionalists out weights the cost.

When a connection is established with the provider application, a package of metrics and data type (all the data does not change during the acquisition as metadata) is sent describing the context of the data collected. The metadata is necessary because different sensors might sample data in different environments, and some applications are depending on recognizing the environment of the data acquisition. Therefore, it is critical to know what data values are measured. Consequently, exposing sensors and data channels through one common interface requires a field of metadata which can be used to: (1) *distinguish* sensor wrapper and data channel the data originated from; (2) determine the *capabilities* of the sensors (i.e., EEG, ECG, LUX); (3) determine the *unit* the data is represented in (i.e., for temperature, Celsius or Fahrenheit); (4) describing the data channel (i.e., placement of the sensor); and (5) a time stamp of when the data was sampled.

To summarize, the task of a *sensor wrapper* is to establish a connection to and collect data from exactly one specified data source, and to send the collected data to the *provider application* that is listening for it. A data source (e.g., BiTalion) can have support for multiple sensor attachments (defined as data channel in the thesis), although, only one sensor wrapper is necessary for each data source and their data channels. Each sensor wrapper is tailored to adapt to the data source's Link Layer technology and the communication protocol of a respective data source. Upon activation by a provider application, the data is collected by the sensor wrapper, and pushed to the provider application in a JSON-format. An illustration of the structure is visualized in Figure 2.1.

2.1.2 Extensible Data Stream Dispatching Tool

The extensible data acquisition tool developed by Gjøby leaves some space for improvements. Such improvements are discussed in the thesis "Extensible data streams dispatching tool for Android" by Daniel Bugajski [9]. Bugajski analyses the potential improvements of the data acquisition tools, which can be extracted into:

- *Lack of reusability*
Only the components that have started the collection can receive the data, and no other components can access the collected data.
- *Lack of sharing*

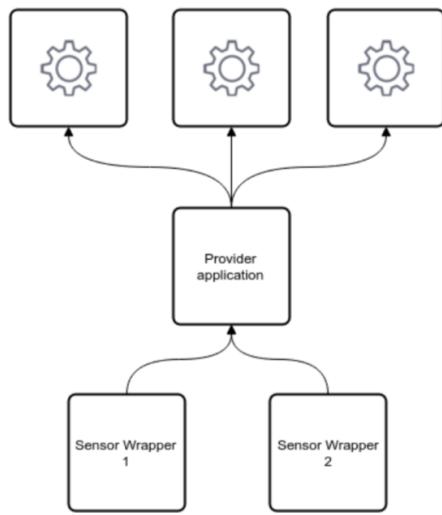


Figure 2.1: Sharing the collected data between multiple applications [22]

Components that perform specific analysis on the collected data in real-time, have no way of share the results of the analysis such that other components can use them.

- *Lack of tuning*

It is not allowed to change the frequency of collection after the start. Thus, the user has to stop the collection and manually change the frequency of the sensor and then restart the collection.

- *Lack of customization*

The set of channels cannot be changed during a collection, and the collector receives data from all channels even if it needs only one of them. Thus, the data packet size and resource usage become larger than necessary.

In the thesis, the modularity of the architecture is improved by extracting the functional requirement of the , and determining the responsibility of each element by. First, finding a model of all available data channels should be implemented. Then, developing a mechanism for cloning a data packet to allow reusing of data across modules. Finally, letting the modules have support for choosing channels they want to receive data from and publish their data to. In the model, these components are distinguished as:

1. *sensor-capability model*: is a representation of all distinct data types and contains all information about the channel. A sensor board usually reads and sends different type of data to a mobile device. Thus, this module is used to control every available data type, such that they can be accessed from the application part at any time.
2. *demultiplexer (DMUX)*: is a data cloner, that receives data packets from one input (i.e., from one channel), and duplicates the data several times - based on the number of subscribers.

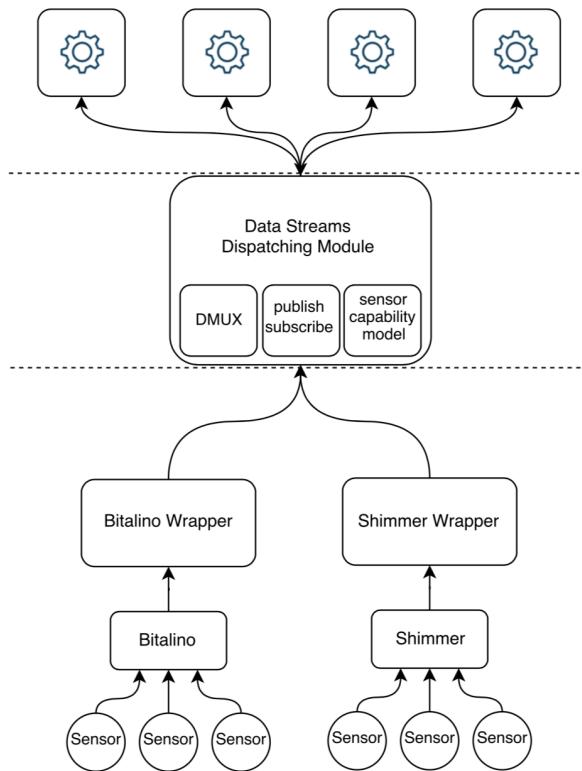


Figure 2.2: Sharing the collected data between multiple applications [9]

3. *publish-subscribe* mechanism: is an interface responsible for providing possibilities of becoming a subscriber or a publisher, in addition, to be able to terminate these statuses. Additionally, every module from the application will be able to see all capabilities represented by this component, enabling the option to choose a frequency the data should be collected with.

The combination of how these modules cooperate and communicate with each other affects the modularity and performance of the architecture. In the thesis, there are various proposed solutions. The naive solution was to fit all of the elements into each respective sensor wrapper, thus, prioritizing the performance and low resource usage, but making it impossible to distinguish data type from two sensor wrapper with the same data type. This is optimal for the cases where collection only occurs on one sensor board. An improved solution includes to place the demux between the application part and the wrapper layer and to insert the remaining elements in the respective application part. This solution resolves the overload of sensor wrappers performing other tasks besides collecting data; thus, the wrappers are untouched, and they send the data to the demux. However, there are several issues with this solution, e.g., due to various obstacles such as (1) every application module has to configure its sensor-capability model; (2) filter requested data packets from all the channels; (3) and deal with the collection speed on its own.

Addressing these issues leads us to the final architecture, which is presented in Figure 2.2, and meets the demands identified in the requirements. In this solution, all elements are placed between the application part and the wrapper layer, forming the data stream dispatching module. The sensor wrapper connects directly to the data streams dispatching module; the module discovers all installed wrappers and populates the sensor-capability model with the data types from all installed sensors. By this, all applications can access a shared sensor-capability mode. A publish-subscribe mechanism enables application modules to subscribe to any capabilities with a preferred sampling rate. Correspondingly, an application module can publish data to other applications through the same interface. The demultiplexing element creates for each subscriber a copy of the data packet.

These three elements together establish *the data stream dispatching module*. The final architecture has a couple of advantages, such as it is exceedingly extensible due to its maintainability. For instance, all communication with other layers occurs through one interface; this way, new instances can be added at any time, without the need for modifying large parts of the system. The system is also efficient due to packets are immediately sent to the application on request (without any buffers), and packets are only sent to the application requesting them, resulting in less resource and power usage, and more battery.

2.1.3 Flow Sensor Kit

Flow is an activity sensor created by SweetZpot, initially designed for measuring breathing and heart rate during activities. The sensor measures breathing per minute, in correlation to the heart rate, for an optimal activity measurement. As stated on their page: *"Usually, at rest, your breathing varies between 6 and 8 liters per minute. During sleep, breathing can be as low as 3 liters per minute and can reach 160 liters per minute and above during high intensity athletic activity"* [39]. Thus, this sensor could be suiting for measuring sleeping problems in the project.

The sensor is a strap-on placed beneath the chest. It weights in at 27 grams, and a dimensions of 77x43x17mm. Additionally, it is equipped with a 3V Lithium battery, with an estimated battery life of one year (with 7 hours a weekly usage). The sensor can be connected with an hand held device through Bluetooth technology, making it possible to connect with a mobile device.

The sensor is an affordable piece of equipment that we can utilize in the project in addition to the sensor we already have (i.e., BiTalino). The product has open-source library to support gathering sensor data, that we eligibility can use to analyze the collected data. Thus, creating a new wrapper for this sensor is required. Consequently, we can combine the sensor data with already integrated sensor in the system, to normalize the

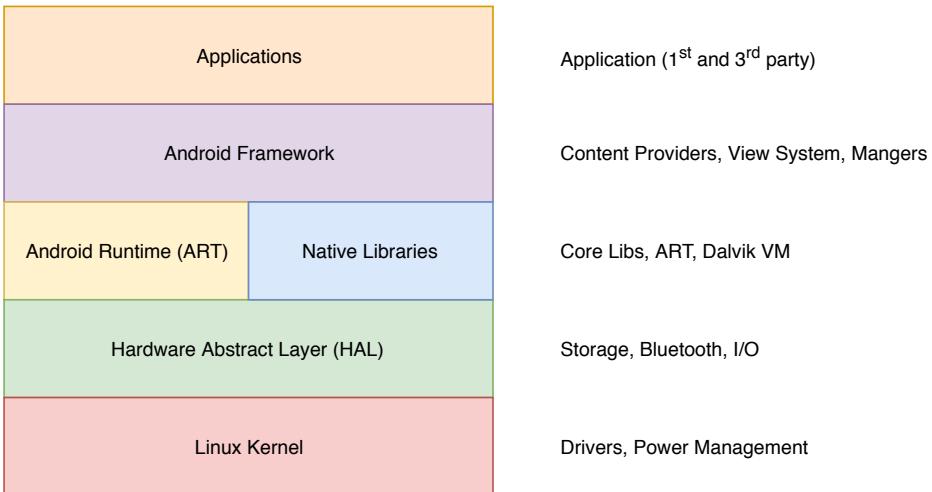


Figure 2.3: Recording

data and minimizing the variance for a more accurate measurement.

2.2 Android OS

Android is an operating system (OS) developed by Google Inc.

2.2.1 Android Architecture

The Android platform is an open-source and Linux-based software stack, containing six major components [34]:

- **Applications:** Android provides a core set of applications (e.g., SMS, Mail, and browser) pre-installed on the device. There are support for installing third-party applications, which allows users to install applications developed by external vendors. A user is not bound to use the pre-installed applications for a service (e.g., SMS), and can choose the desired applications for a service. Also, third-party applications can invoke the functionality of the core applications (e.g., SMS), instead of developing the functionality from scratch.
- **Android Framework:** Is the building blocks to create Android applications by utilizing the core, all exposed through an API. The API enables reuse of core, modular system components, and services; briefly characterized as *View System*: to build the user interface pre-defined components (e.g., lists, grids, and buttons); *Resource Manager*: provides access to resources (e.g., strings, graphics and layout files); *Notification Managers*: allows applications to show custom notifications in the status bar; *Activity Manager*: manages lifecycle of

the application; and *Content Providers*: enables applications to access data from other applications.

- **Android Runtime:** Applications run its own process and has its own instance of the Android Runtime (ART). ART is designed to run on multiple virtual machines by executing DEX (Dalvik Executable format) files, which is a bytecode specifically for Android to optimize memory footprint. Some of the features that ART provides are ahead-of-time (AOT) and just-in-time (JIT) compilation, garbage collection (GC), and debugging support (e.g., sampling profiler, diagnostic exceptions and crash reporting).
- **Native Libraries:** Most of the core Android components and services native code, that requires native libraries, is written in C or C++. The Android platform exposes Java APIs to some of the functionality of the native libraries (with Android NDK).
- **Hardware Abstract Layer:** Provides an interface to expose hardware capabilities to the Java API framework. Hardware Abstract Layer (HAL) consists of multiple library modules that implements an interface for specific hardware components (e.g., camera, or bluetooth module).
- **Linux Kernel:** is the fundation of the Android platform. The ART relies on the functionality from the Linux kernel, such as threading and low-level memory management. The Linux kernel provides drivers to services (e.g., Bluetooth, Wifi, and IPC), and incorporates a component for power management.

2.2.2 Application Components

Application components consists of four core components that are the building blocks of an Android application [6]. This section introduces these components; Activities, Services, BroadcastReceivers, and Content Providers. The activity is responsible for interactions with the user, services is a component that perform (long-running) tasks in the background, broadcast receivers handles broadcast messages from application components, and content providers manages shared set of application data. To enable the components, the Android system must be aware of the components existence. The existence of the components are defined in the manifest file (`AndroidManifest.xml`), which describes the component and the interactions between them, as well as describe the permission of the application.

2.2.2.1 Activity

An application can conists of multiple acitvities, and an activity represents a single screen with a user interface [2]. Applications with multiple

activities has to mark one of the activities as main activity, which will be presented to the user on launch. The user interface of an activitiy is constructed in layout files which define the interaction logic of the user interface, and the layout file is inflated into the activity on launch.

Activites are placed on a stack, and the activity on top of the stack becomes the running activity. Previous activties remain in the stack (unless discarded), and are brought back if desired. An activity can exist in three states cite[activites]:

- **Resumed (Running):** The activity is in the foreground of the screen and has user focus.
- **Paused:** Another activity is running, but the paused activity is still visible. For instance, the other acitvity does not cover the whole screen. A paused activity maintain its state, but can be killed by the system if the memory situtation is critical.
- **Stopped:** The activity is obscured by another activity. A stopped activity maintain its state; however, it is not visible to the user and can be killed if the memory situtation is critical.

Paused and stopped activities can be terminated due to insufficient memory by asking the activity to finish. When the paused or stopped activity is re-opened, it must be created all over.

Activites are part of a activity lifecycle, in Figure 2.4, the state of the activity can be Vaugh categorized into:

- **Entire Lifetime:** of an activity occurs between the calls to `OnCreate()` and the call to `OnDestroy()`. The activity sets the states (e.g., defining the layout) in `OnCreate()`, and release remaning resources in `OnDestroy()`.
- **Visible Lifetime:** of an activity happens between the calls to `onStart()` and the call to `onStop()`. Within this lifecycle, the user can see and interact with the application. Any resources that impact or effect the application occurs between these metods. As activties can alternative between state, the system might call these methods multiple times during the lifecycle of the activity.
- **Foreground Lifetime:** of an activity occurs between the calls to `onResume()` and `onPause()`. The activity is on top of the stack, and has user input focus. An activity can frequently transition in this state; therefore, ensuring that the code in these methods are lightweight in order to prevent the user from waiting.

Fragment

A fragment represents a behavior or is a part of a user interface that can be placed in an activity [19]. Fragment allows for reuse of user interface or behavior across applications, and can be combined to build a multi-pane user interface inside an activity. The fragment allows for more flexibility

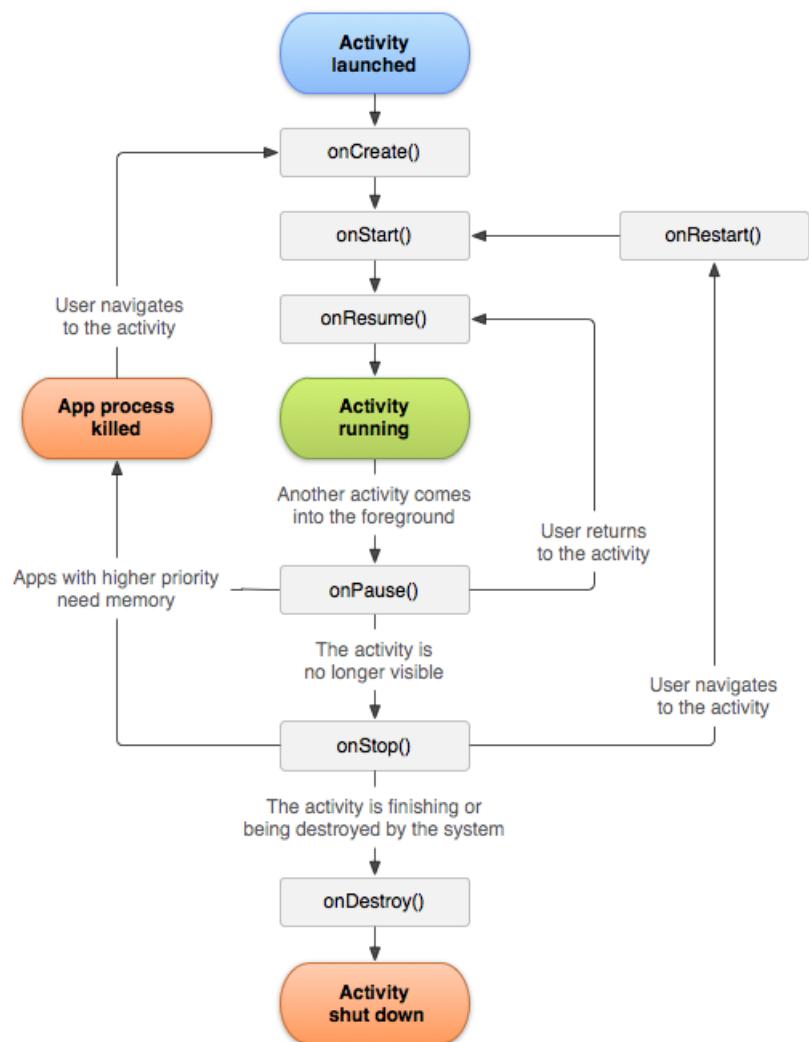


Figure 2.4: Recording

around the user interface, by allowing activities to comprise of multiple fragments which will have their own layout, events, and lifecycles. The lifecycle of a fragment is quite similar to the activity lifecycle; with extended states for: fragment attachment/deattachment, fragment view creation/destruction, and host activity creation/destruction. A fragment is coherent with its host activity, and the state of the fragment is affected by the state of the host activity. Fragment creation and interactions is done through `FragmentManager`.

2.2.2.2 Service

Service is a component that runs in the background to perform long-running tasks [38]. The application or other applications can start a service which remain in the background even if the user switches applications. In contrast, activities are not able to continue if user switches to another application. Also, a service can bind with a component to interact or perform inter-process communication (IPC). To summarize, a service has two forms:

- **Started:** A component can call the `startService()` method on a service, such that the service can run in the background.
- **Bound:** A component can call the `bindService()` method on a service, which in return will offer a client-server interface to perform operations (e.g., sending requests, or retrieving results) across processes with inter-process communication (IPC). Multiple component can bind to a service, and the last component to unbind will destroy the service.

2.2.2.3 Broadcast Receiver

A broadcast receiver is a component that receives broadcast announcements mostly originating from the system (e.g., screen turned off, battery is low, or a picture was captured). Applications can subscribe to messages, and the `BroadcastReceiver` can address and process the messages accordingly. Applications can also initiate broadcasts, and the data is delivered as an `Intent` object. A `BroadcastReceiver` can be registered in the activity of the application (with `IntentFilter`), or inside of the manifest file.

2.2.2.4 Content Providers

Content providers manage access to a set of structured data, and provide mechanism to encapsulate and secure the data [12]. Content providers is an interface which enables one process to connect its data with another process. Also, in order to copy and paste complex data or files between applications, a content provider is required. For instance, to share a file

across a media (e.g., mail), a `FileProvider` (subclass of `ContentProvider`) is needed to facilitate a secure sharing of files [18].

2.2.3 Process and Threads

The Android system creates a Linux process with a single thread of execution for an application on launch [36]. All components (i.e., activity, service, broadcast receiver, and providers) run in the same process and thread (called the *main* thread), unless the developer arrange for components to run in separate process. A process can also have additional threads for processing.

When the memory on the device runs low and demanded by processes which are serving the user, Android might kill low priority processes. Android decides to kill the process based on priority; the process hierarchy consists of five levels (lowest priority number is the most important and is killed last):

1. **Foreground Process:** is a process that is required by the user to interact and function with the application. A foreground process is often an activity that the user interacts with, a service that is bound to an interacting activity, a service that is running in the foreground (with `startForeground()`), a service that is executing on of the lifecycle callbacks, and broadcast receivers executing `onReceive()` method.
2. **Visible Process:** is a process without foreground components, but affect the user interactions. A visible process is when a foreground process takes control (however, the visible process can be seen behind it), and a service that is bound to a visible (or foreground) activity.
3. **Service Process:** is a process that execute work which is not displayed to the user (e.g., playing music or downloading data), and are started with the `startService()` method.
4. **Background Process:** is a process that holds information of paused activities. This process state has no impact on the user experience, and these processes are kept in a LRU (least recently used) in order to refrain killing the activity that the user used last. The state of the process can be saved, if the lifecycle method in activity is implemented correctly, to ensure a seamless user experience.
5. **Empty Process:** is a process that do not hold any active application components; however, are kept alive for caching and faster startup time for components that needs to be executed.

Threads

The main thread is responsible for dispatching events to the user interface widget and drawing events. Also, the thread interacts with the application components from the Android UI toolkit; the main thread is also called

UI thread. System calls to other components are dispatched from the main thread, and components that run in the same process are instantiated from the main thread. Intensive work (such as long-running operations as network access or database queries) in response to user interaction, can lead to blocking of the user interface. As a consequence, the user can find the application to hang, and might decide to quit or uninstall the application.

Additionally, the tool kit to update the user interface in Android is not thread-safe; therefore, enforcing the rules of 1) not to block the UI thread; and 2) not to access the Android UI toolkit outside the UI thread. In order to run long-running or blocking operations, one can spawn a new thread or Android provides several options: `runOnUiThread`, `postDelay`, and `AsyncTask` (perform asynchronous task in a worker thread, and publishes the results on the UI thread).

2.2.4 Inter-Process Communication (IPC)

Inter-process communication is a mechanism to perform remote procedure calls (RPC) to application components that are executed remotely (in another process), with results returned back to the caller. To perform IPC, the caller has to bind to a remote service (using `bindService()`). Upon binding to a remote service, a proxy used for communication with the remote service is returned. The proxy decomposes the method calls, and the Binder framework takes the methods and transfers them to the remote process [7]. Android offers a language to enable IPC; called Android Interface Definition Language (AIDL) [5].

Besides AIDL, one can use Intent to pass messages across processes. Intent is a messaging object used to request action from another application components [24]. There are two types of intents: 1) explicit intents: used to start a component in the application, by supplying the application package name or component class; and 2) implicit intents: declare a general action to perform, which enables other applications to handle it. The main use cases of an intent are to starting an activity; starting a service; and delivering a broadcast.

2.2.5 Data and File Storage

Android provides options to store application data on the device, depending on space requirement, data type, and whether the data should be accessible to other application or private to the application [15]. There are four distinctive data storage options, depending on the requirement of data that is being stored:

- **Internal File Storage:** The system provides a directory on the file system for the application to store and access files. By default, the files

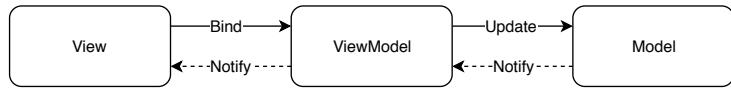


Figure 2.5: Entity Relationship Diagram

saved in this directory is private to the application. Also, files stored in the internal storage are removed on uninstallation of the application; therefore, storing persistent data that are expected to be on the device regardless of the application removal, should not be using internal file storage. In addition, internal file storage allows for caching, which enables temporarily data storage (that do not require to be persistent).

- **External File Storage:** External file storage enables storing of files in a space where users can mount to a computer as an external storage device (or to a physical removable storage, such as SD card). Files stored in external storage makes it possible to transfer files on a computer over USB. Files stored in external file storage enables other applications to access the data, and the data remain available after application uninstallation (unless specified that the storage is application-specific).
- **Shared Preferences:** For storing small and unstructured data, SharedPreference enables API to read and write persistent key-value pairs of primitive data types (e.g., booleans, floats, ints, longs, and strings). The storage location is specified by uniquely identifying name, and the data is stored into a XML file. Also, the data stored remain persistent (even after application termination).
- **Databases:** Android provides support for SQLite databases, which is a relational database management system embedded into the system. The access to the database is private to the application, and accessing the database can be done with the Room persistence library. The Room library provides an abstract layer over SQLite APIs.

2.2.6 Architecture Patterns

The architectural pattern principle enhances the separation of *graphical user interface* logic from the *operating system* interactions [23]. The Model-View-ViewModel (hereafter: MVVM) is an architectural pattern which is well-integrated and incentivized by Android. It has three components that constitute the principle:

- **Model:** represents the data and the business logic of the application.
- **ViewModel:** interacts with the model, and manages the state of the view.
- **View:** handles and manages the user interface of the application.

In Figure 2.5, the interactions amongst the components are illustrated. The connection between the View and ViewModel occurs over a data binding connection, which enables the view to change automatically based on changes to the binding of the subscribed data [29]. In Android, the LiveData is an observable data holder that enables data binding, which allows components to observe for data changes. LiveData respects the lifecycle of the application components (e.g., activities, fragments, or services), ensuring the LiveData only updates the components that are in an active lifecycle state [28]. Moreover, Android Room provides set of components to facilitate the structure of the model component [37]. More specifically, it models a database and the entities (which are the tables in the database).

2.2.7 Power Management

Battery life is a concern to the user, as the battery capacity is significantly limited on devices [35]. Android has features to extend battery life by optimizing the behavior of the application and the device, and provides several techniques to improve battery life:

- **App Restrictions:** Users can choose to restrict applications (e.g., applications cannot use network in the background, and foreground services). The application that are restricted, functions as expected when the user launches the application; however, are restricted to do background tasks.
- **App Standby:** An application can be put into standby mode if an user is not actively using it, resulting in the application background activity and jobs is postponed. An application is in standby mode if there is no process in the foreground, a notification is being viewed by the user, and not explicitly being used by the user.
- **Doze:** When a device is unused for a long period, applications are delayed to do background activities and tasks. The doze-mode enters maintenance window to complete pending work, and then resume sleep for a longer period of time. This cycles through until a maximum of sleep time is reached. Some applications want to keep the device running to perform long-running tasks (e.g., collecting data), and WakeLocks enables this. WakeLocks allows application to perform activities and tasks, even while the screen is turned off [25].
- **Exemptions:** Another way of keeping an application awake, is to exempting applications from being forced to Doze or to be in App Standby. The exempted applications are listed in the settings of the device, and users can manually choose application to exempt. Consequently, exempted applications might overconsume the battery of the device.

2.2.8 Bluetooth Low Energy

Android supports for Bluetooth Low Energy (BLE), which is designed to provide lower power consumption on data transmission (in contrast to classic Bluetooth) [8]. BLE allows Android applications to communicate with sensors or devices (e.g., heart rate sensor, and fitness devices), that has a stricter power requirements. Sensors that utilize BLE, are designed to last for a longer period of time (e.g., weeks or months before needing to charge or replace battery). The protocol of BLE is optimized for small burst of data exchange, and terms and concepts that form a BLE can be characterized as:

- **Generic Attribute Profile (GATT):** As of now, all Low Energy applications are based on GATT. GATT is a general specification for sending and receiving burst of data (known as *attributes*) over a BLE link.
- **Attribute Protocol (ATT):** GATT utilizes the Attribute Protocol, which uses a few bytes as possible to be uniquely identified by a Universal Unique Identifier (UUID). An UUID is a standardized format for identifying information.
- **Characteristic:** Contains a single value and descriptors that describe the characteristic value (i.e., can be seen as a type).
- **Descriptor:** Are defined attributes that describe a characteristic value (e.g., specify a human-readable description, a range of acceptable values, or a unit of measure).
- **Service:** Is a collection of characteristics (e.g., a service is *Heart Rate Monitor* which includes a characteristic of *heart rate measurement*).

In order to enable BLE, facilitating for a GATT server and GATT client is required. Either the sensor or the device take the role of being a server or a client. However, the GATT server offers a set of services (i.e., features), where each service has a set of characteristics. And the GATT client can subscribe and read from the services the GATT server provides.

Chapter 3

Related Work

The widespread adoption of detecting and analyzing sleep-related breathing disorders on a mobile device has been a research topic and concert for some time. Various techniques and methods has been applied to detect sleep-related breathing disorders on a mobile phone with the use of built-in or external sensors. We will survey some of the research conducted in this field, and in the end, we will present similiarites based on our system.

Nandakumar, Gollakota, and Watson [32] presents a contactless solution for detecting sleep apnea events on smartphones. The goal behind the research is to detect sleep apnea events without any sensor on the human body. They achieved this by transforming the phone into an active sonar system, that emits frequency-modulated sound signals and observe for the reflections. Based on the experiments, the system operates efficiently at distances of up to a meter, also while the subject is under a blanket. They performed a clincal study with 37 patients, and concluded that the system managed to accurately compute the central, obstructive, and hyponea events (with a correlation coefficient of 0.99, 0.98, and 0.95).

Alqassim *et al.* [3] designed and implemented a mobile application (for Windows and Android) to monitor and detect symptoms of sleep apnea using built-in sensors in the smart phone. The purpose of the application is to make users aware of whether they have sleep apnea, before they continue with a more expensive and advanced sleep tests. They achieved this by measuring the breathing patterns and movements patterns based on the built-in microphone and accelerometer in the smart phone. The system instructs the user to place the smart phone on its arm, abdomen or near the bed during recording. The data is collected on the smart phone, and sent to a centeral server in the cloud, where authroized docots can review the samples. To summarize, the system tries to monitor sleep apnea in the aspect of motion and voice recorder, in order to detect sleep apnea on a smart phone.

Penzel *et al.* [33] investigates the challenges and develop a system assoicated with insufficient conventional sleep laboratoris and their expensive

and time-consuming polysomnographic diagnostics as early as in 1989. The purpose of the system is to make diagnose of sleep-related breathing disorders available at any hospital; the system was placed in a wooden case with wheels to be moved between bedside locations. They developed a circuit board that has support for various sensor to evaluate breathing, ECG, blood gases and the state of sleep. During record, the samples from the sensors were compressed to a resolution of one value per second per parameter (e.g., mean value of respiratory frequency, ventilated volume, actigraph activity, and EOG activity) and stored on a personal computer. After the recording, manual evaluation can be carried out using print-outs, as well as reviewing the data on a screen. Thus, the software of the system facilitates recording and reviewing of the data, also basic evaluation and analysis.

To summarize, Nandakumar, Gollakota, and Watson created an application to collect samples through a contactless solution, which quite accurately measures central, obstructive, and hypopnea events. Alqassim *et al.* developed a mobile application to sample breathing patterns and movement patterns with the use of the built-in microphone and accelerometer in the smartphones. And Penzel *et al.* built a mobile system to record and analyze sleep-related breathing disorders with technology that was advanced at the time.

To conclude this chapter, we can derive that the demand for a mobile system to detect sleep-related breathing disorders is high. Developing a system that records, analyzes and detects sleep-related breathing disorders (e.g., Obstructive Sleep Apnea) enables users to diagnose the sleeping disorder from home, in contrast to a expensive and time-consuming polysomnographic diagnostics at a laboratory. Therefore, developing a system that is extensible and modular to new techniques and methods, as well as for future sensors devices, is appropriate in order to extend and evolve the research behind the detecting sleep-related breathing disorders on a mobile device. We developed an application for Android which enables recording, sharing and analytics, in addition to support for third-party modules to extend the functionality of the application, and to enrich the data collected on a patient's device. The application is called *Nidra*, which we present in the next chapter.

Part II

Design and Implementation

Chapter 4

Analysis and High-Level Design

It is the goal of this thesis to enable detection of sleep-related illnesses with the aid of an Android device and low-cost sensors, and to further analyze and evaluate sleep- and breath-related patterns. We developed an application, called *Nidra*, which attempts to collect, analyze and share data collected from external sensors, all on a mobile device. Also, Nidra acts as a platform for modules to enrich the data, thus extending the functionality of the application.

The motivation behind this application is to provide an interface for patients to potentially run a self-diagnostic from home, and to aid researchers and doctors with analysis of sleep- and breathing-related illnesses (e.g., Obstructive Sleep Apnea). An overview of the Nidra application pipeline can be found in Figure x, beginning with data acquired from a sensor, and ending with the data in the Nidra application. As for now, Nidra consists of three main functionalities, each related to the requirements defined in Section [Problem Statement].

1. The application should provide an interface for the patient to 1) record physiological signals (e.g., during sleep); 2) present the results; and 3) share the results.
2. The application should provide an interface for the developers to create modules to enrich the data from records or extend the functionality of the application.
3. The application should ensure a seamless and continuous data stream, uninterrupted from sensor disconnections and human disruptions.

This chapter will give a detailed look at the design of Nidra, including the tasks which constitute the structure of the application, the separate concerns in relation to the tasks, and the structure of the data in the application.

4.1 Requirement Analysis

4.1.1 Stakeholders

McGrath and Whitty [31] describe the term stakeholder as those persons or organizations that have, or claim an interest in the project. They distinguish stakeholders into four categories: 1) *contributing (primary) stakeholders* participate in developing and sustaining the project; 2) *observer (secondary) stakeholders* affect or influence the project; 3) *end-users (tertiary stakeholder)* interact and uses the output of the application; and 4) *invested stakeholders* has control of the project. In Nidra, there are three stakeholders who affect the application, and each can be categorized respectfully:

- Patients: Are identified as an end-user; they interact with the application.
- Researchers/Doctors: Are identified as an observer stakeholder; they might not use the application itself; however, they might use the data obtained from the patients' recordings for further analysis. Also, request functionality in the application.
- Developers: Are identified as a contributor stakeholder; they maintain the application from bugs or extend the functionality of the application. Additionally, they can contribute to developing modules that extend the functionality of the application.

4.1.2 Resource Efficiency

The application is designed for the use on a mobile device; modern mobile devices are empowered with multi-core processors, a sufficient amount of ROM, and a variety of sensors. However, the battery capacity is restrictive and based on usage. The device may only last for one day before a charge, due to the size of the battery capacity [27]. The average battery capacity of a mobile device is approximately 2000 mAh on budget devices and around 3000 mAh on high-end devices [17]. The application should be able to run at least 7 hours without any power supply. Also, the device should be capable of handling various sensor connections simultaneously. Therefore, the application should be designed to be resource efficient, by utilizing least amount of battery resources during a recording. Also, ensure sufficient amount of power on the device before starting a recording session.

4.1.3 Security and Privacy

The proposed use of the application is to monitor the sleeping patterns of a patient. The application manages and stores personal- and health-related data about the patient. As a precaution, the application should incorporate the CIA triad, which stresses data confidentiality, integrity, and availability

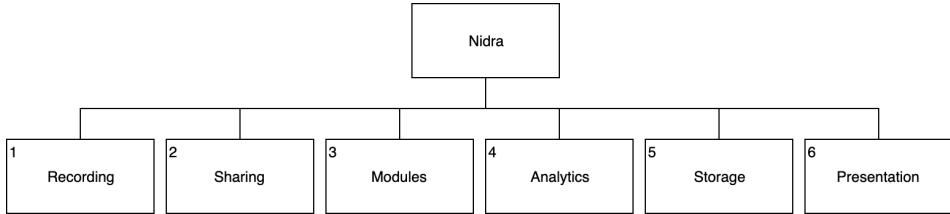


Figure 4.1: Recording

[42]. Any unauthorized access to the data, data leaks, and confidentiality should be appropriately managed on the device. Sharing the data across application or with researchers/doctors should be granted with the consent from the patient. Besides, a mobile device can be connected to the Internet, which makes it vulnerable to attacks. Also, other installed application on the device can manipulate the access to the data. Therefore, revising the security policy defined by Android [4] should be incentivized.

4.2 High-Level Design

4.2.1 Task Analysis

Task analysis is a methodology to facilitate the design of complex systems. Hierarchical task analysis (HTA) is an underlying technique that analyzes and decomposes complex tasks such as planning, diagnosis, and decision making into specific subtasks [13]. In Figure 4.1, an illustration of the the tasks of the application is presented. This Subsection will introduce these tasks, which are an integral part to the development of the application.

Recording

A *recording* is a process of collecting and storing physiological signals from sensors over an extended period (e.g., overnight). To enable a recording, we need to establish connections to available sensors, collect samples from the sensors, and storing the samples on the device. A *sensor* is a device that transforms analog signals from the real world into digital signals. The digital signals are transmittable over Link Layer technologies (e.g., Bluetooth), and the communication between a sensor and device occurs over an application programming interface (API). A *sample* is a single sensor reading containing data and metadata, such as time and the physiological data. During a recording session, ensuring for a consistent and uninterrupted data stream from the sensors is vital to obtaining persistent and meaningful data. Once a recording session has terminated, a *record* with metadata about the recording session is stored, alongside the samples.

Sharing

Sharing is a mechanism to export and import records across applications.

Exporting consists of bundling one or more records with correlated samples into a transmittable format, and transferring the bundled records over a media (e.g., mail). *Importing*, on the other hand, consists of locating the bundled records on the device, parsing the content and storing it on the device. [skrive mer?]

Module

A *module* is an independent application that can be installed and launched in Nidra (hereafter: application), to provide extended functionality and data enrichment. A module does not necessarily interact with the application; however, it utilizes the data (e.g., records). For example, a module could be using the records to feed a machine learning algorithm to predict obstructive sleep apnea. Installing a module is achieved by locating the module-application on the device, and storing the reference in the application. Due to limitations in Android, the module-application cannot be executed within the application. Therefore, the module-application is a standalone Android application; furthermore, the development of the module-application is independent from the application.

Analytics

Analytics is the visualization and interpretation of patterns in the records. The application facilitates the recording of physiological signals, which enables the detection and analysis of sleep-related illnesses. There are various analytical methods, ranging from graphs to advanced machine learning algorithms. Incorporating a simple time series plot can indirectly aid in the analysis. For instance, plotting a time series graph where the physiological signals are on the Y-axis and the time on X-axis, provides a graphical representation of the data that can be further analyzed within the application.

Storage

Storage is the objective of achieving persistent data; data remain available after application termination. To enable storage, we use a database for a collection of related data that is easily accessed, managed, and updated. The database should be able to store records, samples, modules, and biometrical data related to user (i.e., gender, age, height, and weight). Structuring a database that is reliable, efficient, and secure is a crucial part of achieving persistent storage. Android provides several options to enable storage on the device (e.g., internal storage and database).

Presentation

Presentation is the concept of exhibiting the functionality of the application to the user. A user interface (UI) is the part of the system that facilitates interaction between the user and the system. In Nidra, determining the layout and view of the application, color palette, interactions, and feedback on actions is part of the development of a user interface.

4.3 Separation of Concerns

Separation of concern is a paradigm that classifies an application into concerns at a conceptual and implementational level. It is beneficial for reducing complexity, improving understandability, and increasing reusability [26]. The concerns in this thesis are the individual tasks defined in task analysis. Each concern is conceptualized with a graph of components, the functionality of each component when combined constitutes a structure. The structure of each concern is derived based on research and development. In this Section, we will analyze and decompose the tasks defined in task analysis into subtasks (hereafter: components), where each component is functionality

4.3.1 Recording

The structure of a recording is restrictive in terms of arranging the components due to the design of CESAR. There are numerous ways of presenting the recording view; however, a recording structure is limited to the components of starting a recording, establishing sensor connection, monitoring of samples, and finalize sensors and recording. Additional components can be incorporated to aid a recording without causing disruption. For instance, the connectivity state component (see Section 4.3.1.1) provides extended functionality to the recording structure. In Figure 4.2, the illustration of a recording structure with the components and their dependencies are shown:

- 1.1 Sensor Discover: Find all eligible sensors that can enable a recording.
 - 1.1.1 Select Sensors: From the sensor discovery, we can choose preferable sensors sources.
 - 1.1.2 All Sensors: more Straightforward, we sample from all of the available sensors.
- 1.2 Sensor Initialization: Once we have a list of sensors sources, we need to establish and initialize a connection with the sensors. Occasionally a sensor might use some time to connect, or unforeseen occurrence is hindering the initialization of the sensor. Therefore, halting the sensor initialization or actively checking for sensor initialization is important.
- 1.3 Sensor Connectivity Setup: Establish a connection between the application and the sensor source through an API or IPC connection. All data exchange will occur over the established interface.
- 1.4 Connection State: Based on sensors establishments we can proceed to either start or stop a recording.

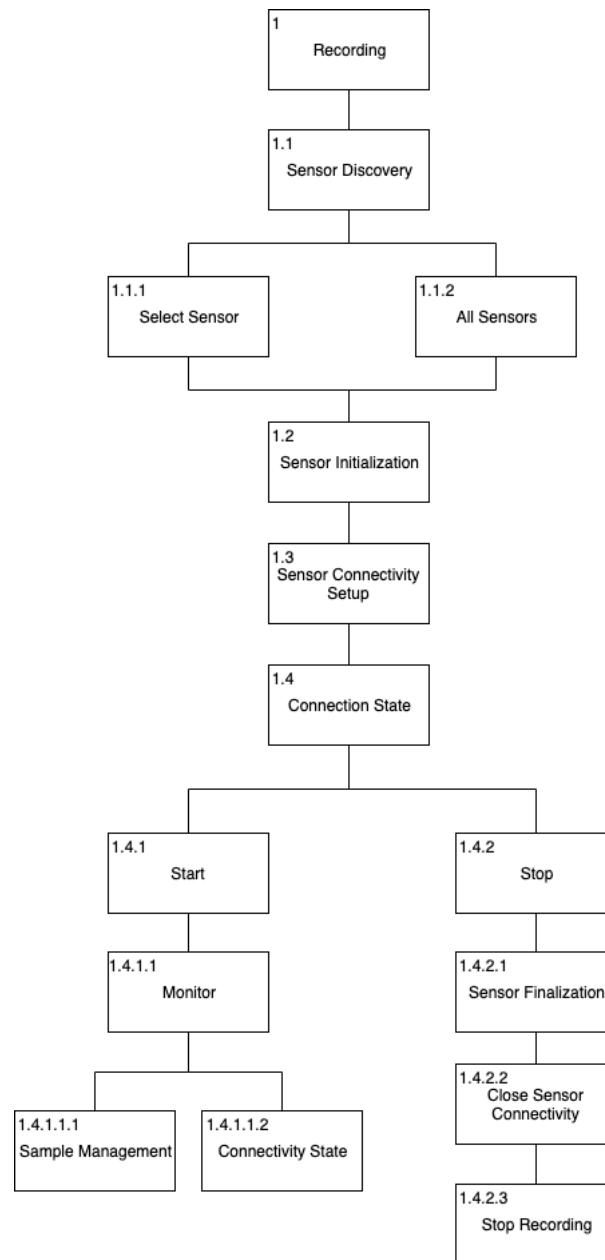


Figure 4.2: Recording

1.4.1 Start: By starting, we notify the sensors to begin collecting data, and the application should display that a recording has begun accordingly. Also, start a timer to display time elapsed on the current recording.

1.4.1.1 Monitor: Is a mechanism to handle the connectivity state and the incoming samples. It is actively listening to the interface for new data from the sensors, and appropriately distributing the data to the sample management component.

1.4.1.1.1 Sample Management: Handles a single sample from a sensor by parsing the content according to the payload of the sensor (each sensor might have different payload structure), such that it is according to our data structure.

1.4.1.1.2 Connectivity State: Actively checking the state of sensor connectivity (read more below)

1.4.2 Stop: Stop the recording timer and proceed to display results.

1.4.2.1 Sensor Finalization: Notify the sensor to stop sampling data, and close establishment.

1.4.2.2 Close Sensor Connectivity: Close the interface establishment between the application and the sensors.

1.4.2.3 Stop Recording: Once the sensors has closed its connections, add additional information to the recording (e.g., title, description, rating). In the end, the recording has concluded and it is stored on the mobile device.

4.3.1.1 Connectivity State Component

Connectivity state is a component that monitors for unexpected sensor disconnections or disruptions. Unexpected behavior can occur due to anomalies in the sensor, or the sensor being out of reach from the device for a brief moment. A naive solution would be to ignore the connectivity state component, and assuming the sensors are connected to the device indefinitely. However, upon disconnections or disruptions, the recording would be missing samples which will result in a lacking record. This component solves the issue of missing samples by actively reconnecting the sensor based on a time interval, resulting in more accurate record with fewer gaps between samples. The following design questions for this component are 1) should the connectivity state component, which implements a time interval that tries to reconnect with the sensor, be implemented in the sensor wrapper, or should it be in the proposed recording structure?; and 2) should the interval between sample arrival be a fixed time or a dynamical time?

1. To achieve a mechanism of reconnecting to the sensor on unexpected disconnects or disruptions, establishing a time interval that monitors

for sample arrivals within a time frame (e.g., every 10 seconds) is required. Incorporating the time interval in the sensor wrapper reduces the complexity of Nidra. However, it introduces extra complexity to the sensor wrappers. A sensor wrapper has to distinguish actual disconnects from unexpected disconnects. Although, by extending the functionality of sensor wrapper by implementing a state that indicates whether a recording is undergoing or stopped solves the problem. All future sensor wrappers would then have to implement the proposed solution, resulting in a complicated and time-consuming sensor wrappers development. While implementing the proposed solution in the sensor wrappers is possible, extending the recording structure with the logic in Nidra would be more meaningful and time-saving. In our design, we will be implementing the connectivity state in the recording structure.

2. A time interval triggers an event every specified time frame. If an event is triggered, a sample has not arrived, meaning the sensor either has been disconnected or disrupted. A time frame can be in a fixed size (e.g., every 10 seconds) or a dynamical size (e.g., start with 10 seconds, then incrementally increase the frame by X seconds). Implementing a fixed time frame increases the stress on put on the sensor, whereas a dynamical time might miss samples if the time frame is significant. Depending on how critical the recording is, a suitable solution for the time frame should be configurable. Also, limiting the number of attempts made to reconnect should be considered, due to actively reconnecting to a sensor that is dead or completely out of reach can cause stress on the device. Thus, stopping the recording once a limited number of attempts has been reached. In our design, we implemented a dynamical time and limited the number of attempts to 10.

4.3.2 Sharing

Sharing is separable into two concerns: export and import. The scope of exporting in Nidra is to select desired records, format and bundle the records into a transmittable file, and distributing the bundle over a media (e.g., mail). The scope of importing is to locate the file on the device, parse the content based on the format, and store it on the device. In Figure 4.3, the structure of sharing is presented with components and their dependencies:

- 2.1 **Import:** Is a mechanism that locates a file, parse the data, and stores it on the device.
 - 2.1.1 **Locate File:** To enable this, the user has to download the file on the device. Then, locate the file on the device by using an interface to browse downloaded files. An interface can be developed; however, using the Android document picker (ref) is more straightforward solution.

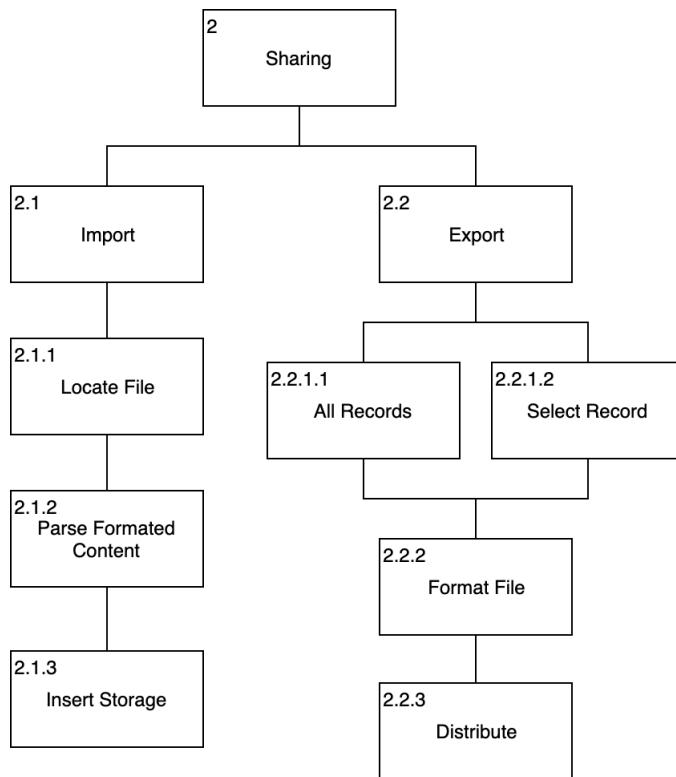


Figure 4.3: Sharing

2.1.2 Parse Formated Content: Parse the content of the file accordingly to the data format discussed in Section 4.4.1.

2.1.3 Insert Storage: Retrieve the necessary data from the parsed file, to store on the device without overriding existing data.

2.2 Export: Is a mechanism that selects all or a specific record, format the record into a formatted (see Section 4.4.1) file, and exports the file across device.

2.2.1.1 All Records: Export all of the records on the device.

2.2.1.2 Select Record: Pick one specific record to export.

2.2.2 Format File: When a preferred format for the records is selected, bundling the data into a formatted (see Section 4.4.1) file for transmittal can be done. It is essential to identify the name of the file uniquely to prevent duplicates and overrides of data. For instance, identifying the name of the file with the device identification appended with the time of exporting.

2.2.3 Distribute: Send the file across application (read more below).

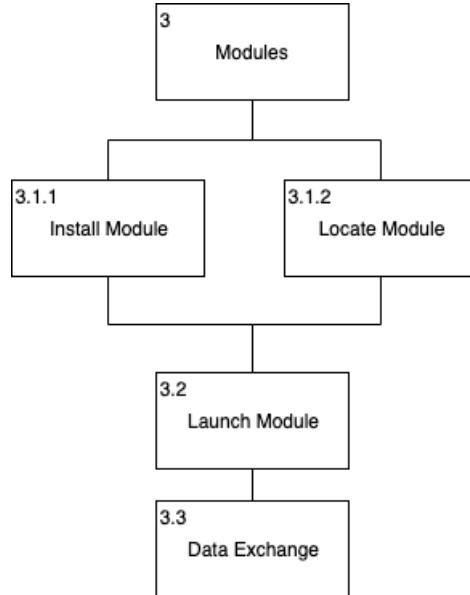


Figure 4.4: Modules

4.3.2.1 Distribute Component

The distribute component uses the formatted file and transfers it across applications. There are two distinctive methods to perform this task which is efficient and practical: 1) implement an interface with recipients to share data with, by establish a web-server with logic to handle users and sharing data with the desired recipient. Also, implementing an interface to retrieve the file within the application; and 2) using the interface provided by Android to share files.

While the first option might be favorable in terms of practicality, this solution introduces additional concerns (e.g., the privacy matters of storing user data on a server) which is out of the scope for this project. For this reason, using the interface provided by Android is a reasonable solution. The user of the application can utilize the Android interface for sharing files over installed applications; however, e-mail is a flexible media to transfer the file, and the user can specify the recipients accordingly.

4.3.3 Modules

Modules are independent applications that provide extended functionality and data enrichment to Nidra. The components for locating and launching a module is limited to Android design; however, the component for data exchange between a module and Nidra can be designed variously. In Figure 4.4, the structure of modules is presented with components and their dependencies:

3.1.1 Install Modules: Is the process of locating the application on the

device, and storing the reference of the application package name in the storage.

- 3.1.2 **Locate Module:** Retrieve the list of stored modules, and display the installed modules to the user.
- 3.2 **Launch Module:** Get the application location stored in the module, and launch the application with the use of Android Intent.
- 3.3 **Data Exchange:** Enrich the module with data from the application (read more below)

4.3.3.1 Data Exchange Component

The data exchange component facilitates the transportation of data between Nidra and a module. As of now, the data is records and corresponding samples, which is formatted (Section Data) accordingly. The two distinct methods to exchange data between a module and Nidra are 1) formatting all of the data and bundling it into the launch of the module, and 2) establishing a communication link for bi-directional requests between Nidra and the module.

Android provides an interface to attach extra data on activity launch. The first solution is, therefore, convenient and efficient; all of the data is formatted and bundled into the launch. However, once Nidra has launched the module-application, there are no ways of transmitting new data besides relaunching the module-application. For this reason, the second option allows for continuous data flow by establishing a communication link with IPC between the applications. The data exchange between Nidra and modules can then be bidirectional; the module can request desired data any time, and Nidra can collect reports and results generated by the module.

One could argue that new records are not obtained while managing and using a module. However, there might be future modules that do a real-time analysis of a recording, but that will require an interface for continuous data flow. For the simplicity of our design, we will be going with the first option of bundling all of the data and sending it on launch.

4.3.4 Analytics

Analytics uses techniques and methods to gain valuable knowledge from data. Nidra provides a simple illustration of the data in a time-series plot; however, other techniques can be incorporated. In essence, the facilitation of modules in the application enables the development opportunities for advanced analytics of the data. In Figure 4.5, the structure of analytics is presented with components and their dependencies:

- 4.1.1 **Select Record:** locate the file on the device.

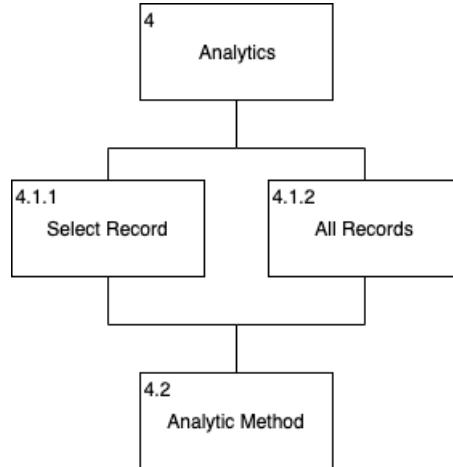


Figure 4.5: Analytics

4.1.2 All Records: parse the content of the file

4.2 Analytic Method:

4.3.4.1 Analytic Method Component

The analytic method component uses the records on the device for representation or analysis. Graphical and non-graphical are two techniques for representing data. Graphical techniques visualize the data in a graph that enables analysis in various ways. A few graphical techniques are diagrams, charts, and time series. Non-graphical techniques, better known as statistical data tables, represent the data in tabular format. This provides a measurement of two or more values at a time [21]. More advanced techniques to analyze the data, are to use machine learning. Machine learning is concerned with developing data-driven algorithms, which can learn from observations without explicit instructions. For example, using recurrent neural networks (e.g., RNN, LSTM) or regression models (e.g., ARIMA), can be used to predict the sleeping patterns [20].

In Nidra, a time series graph is used to represent the data of a record. The time series graph represents the respiration data on the Y-axis and the time on the X-axis. Essentially, the facilitation of modules in the application is designed to enable advanced techniques to predict, analyze, and interpret the data acquisition. Therefore, in Nidra, the analytic methods are limited; however, the modules enable developers to construct any method they desire.

4.3.5 Storage

Storage is the objective of achieving persistent data; data that is available after application termination. The data is characterized into four data

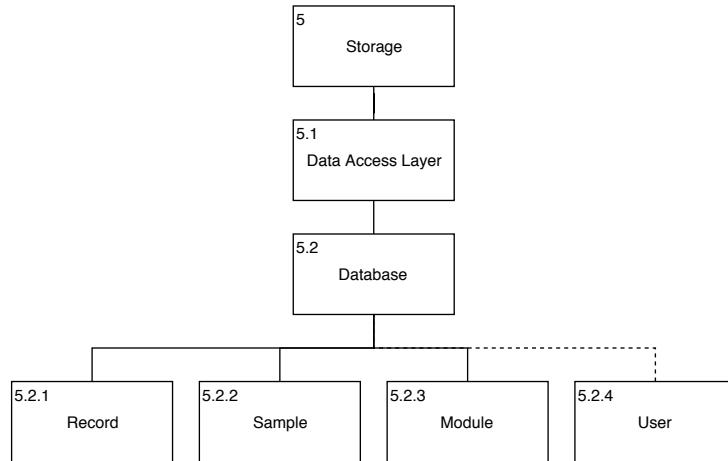


Figure 4.6: Storage

entities (i.e., record, sample, module, and user) that contains individual properties. The components in the storage structure are constructed to be extensible and scalable in terms of future data, restructure of data, and removal of data. In Figure 4.6, the structure of storage is presented with components and their dependencies:

- 5.1 Data Access Layer: Also known as an Data Access Object (DAO), that provides an abstract interface to a database. It exposes spesific operations (such as insertion of a record) without revealing the database logic. The advantage of this interface is to have a single entry point for each database operation, and to easily extend and modify the operation for future data.
- 5.2 Database: Is the storage of all the data (read below).
- 5.2.1 Record: Is a table in the database, which contains fields appropriately to the record structure. A record contains meta-data about a recording (e.g., name, recording time, user). The design decision and an example of a recording record is illustrated in Section 4.4.2.1.
- 5.2.2 Sample: Is a table in the database, which contains fields appropriately to a single sample from a sensor. A sample contains data received from a sensor during a recording. The design decision and an example of a sample record is illustrated in Section 4.4.2.2.
- 5.2.3 Module: Is a table in the database, which contains fields appropriately to the sample structure. A sample contains the name of the module and a reference to the application package. The design decision and an example of a module record is illustrated in Section 4.4.2.3.
- 5.2.4 User: Is an object stored on the device, which contains fields appropriately to the user structure. A user contains the patients biometrical information (e.g., name, weight, height). The design decision and an example of a user is illustrated in Section 4.4.2.4.

4.3.5.1 Database Component

Android provides several options to store data on the device; depending on space requirement, type of data that needs to be stored, and whether the data should be private or accessible to other applications. Two suitable options for storage are 1) internal file storage - storing files to the internal storage private to the application; and 2) database - Android provides full support for SQLite databases, and the database access is private to the application [16]. Based on the options, some following design questions are 1) should the data be stored in a flat file database on the internal file storage, or should it be stored in an SQLite database?

Flat files database encode a database model (e.g., table) as a collection of records with no structured relationship, into a plain text or binary file. For instance, each line of text holds on a record of data, and the fields are separable by delimiters (e.g., comma or tabs). Another possibility is to encode the data in a preferable data format (see Section 4.4.1). Flat file databases are easy to use and suited for small scale use; however, they provide no type of security, there is redundancy in the data, and integrity problems [30]. Locating a record is made possible by loading the file, and systematically iterating until the desired record is found. Similarly, updating a record and deleting a record. Consequently, the design of flat file databases is for simple and limited storage.

SQLite is a relational database management system, which is embedded and supported in Android. Relational database management system (RDBMS) provides data storage in fields and record, represented as columns (fields) and rows (records), in a table. The advantage is the ability to index records, relations between data stored in tables, and support querying of complex data with a query language (e.g., SQL). Also, RDBMS provides data integrity between transactions, improved security, and backup and recovery controls [30].

While a flat file database is applicable to store small and unchangeable data, it is not suitable for scalable and invasive data change. In Nidra, the samples acquisition makes it unreasonable to use a flat file database. Therefore, SQLite is a preferable solution to storing samples. Also, establishing a relationship between a record and a sample is made possible [rewrite].

4.3.6 Presentation

Presentation facilitates the user interface of the application, in terms of visualizing the functionality of the application to the user. The user interface derives from the functionality (concerns discussed) in the application, and research on the topic. In Figure 4.7, the structure of storage is presented with components and their dependencies:

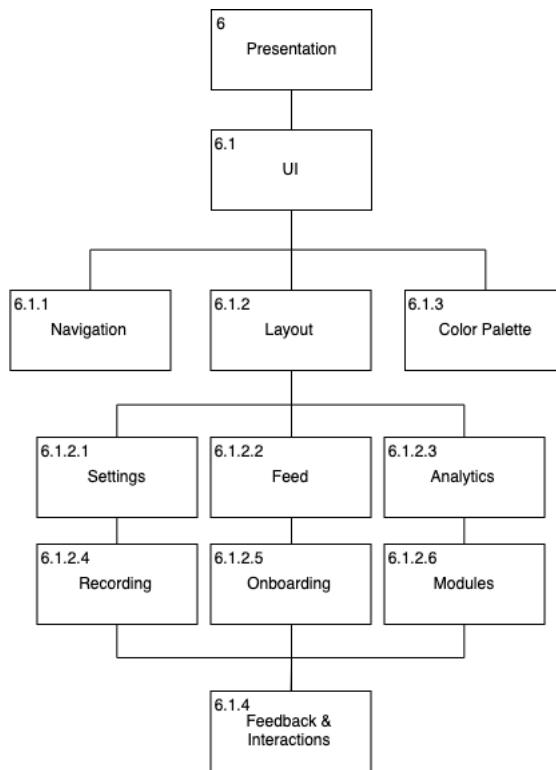


Figure 4.7: Presentation

6.1 UI: The user interface where interaction between users and the application occurs.

6.1.1 Layout: Are the screen with the content of the current screen. The layout incorporates the color palette and some views have the navigation displayed.

6.1.2 Navigation: The navigation is a menu with options to change the layout.

6.1.3 Color Palette: A color selection that persist throughout the application (read more below).

6.1.1.1 Settings: Is a screen with user details, permissions and credits, with options to modify permission and user details.

6.1.1.2 Feed: Is a list of all records for the user, displayed with details specific to the record to make it distinguishable and easily recognizable.

6.1.1.3 Analytics: An interactive time-series graph for a single record.

6.1.1.4 Recording: The process of establishing a recording session, in addition to showing the results after a recording session has ended.

6.1.1.5 Onboarding: The initial screen displayed to the user, where the user can supply the application with their biometrical data.

6.1.1.6 **Modules:** A list of all installed modules, also an option to add more modules.

6.1.4 **Feedback & Interactions:** Each layout has different feedback and interaction, which should be handled appropriately.

4.3.6.1 Color Palette Component

Color palette is a component that decides the color scheme in the application. In the proposal of a color system in the design guidelines by Google [40], it is essential to pick colors that reflect the style of the application accordingly to: 1) primary colors - the most frequently displayed color in the application; 2) secondary colors - provides an accent and distinguish color in the application; and 3) surface, background, error, typography and iconography colors - colors that reflect the primary and secondary color choices.

Moreover, choosing colors that meet the purpose of the application is critical. Nidra is most likely to be used during the evening and the morning. According to Google [14], a dark color theme reduces luminance emitted by the device screen, which reduces eye strain, while still meeting the minimum color contrast ratios, and conserving battery power. From this, we will be choosing a dark color theme.

4.4 Data Structure

4.4.1 Data Formats

The data format is a part of the process of serialization, which enables data storage in a file, transmittal over the Internet, and reconstruction in a different environment. Serialization is the process of converting the state of an object into a stream of bytes, which later can be deserialized by rebuilding the stream of bytes to the original object. There are several data serialization formats; however, JavaScript Object Notation (JSON) and eXtensible Markup Language (XML) are the two most common data serialization formats. In this Section, we will discuss these formats. In the end, we will compare them and choose the format that meets the criteria of being compact, human-readable, and universal.

4.4.1.1 JSON

JSON or JavaScript Object Notation is a light-weight and human-readable format that is commonly used for interchanging data on the web. The format is a text-based solution where the data structure is built on two structures: a collection of name-value pairs (known as objects) and ordered

list of values (known as arrays). The JSON format is language-independent and the data structure universally recognized [1, 43]. However, it is limited to a few predefined data types (i.e., string, number, boolean, object, array, and null), and extending the data type has to be done with the preliminary types.

```
1 {  
2   "user": {  
3     "firstname": "Ola"  
4     "lastname": "Nordmann"  
5   }  
6 }
```

Listing 4.1: My Caption

4.4.1.2 XML

XML or eXtensible Markup Language is a simple and flexible format derived from Standard Generalized Markup Language (SGML), developed by the XML Working Group under the World Wide Web Consortium (W3C). An XML document consists of markups called tags, which are containers that describe and organize the enclosed data. The tag starts with < and ends with >; the content is placed between an opening tag and a closing tag (see listing). [11, 43] XML provides mechanisms to define custom data types, using existing data types as a starting point, making it extensible for future data.

```
1 <user>  
2   <firstname>Ola</firstname>  
3   <lastname>Nordmann</lastname>  
4 </user>
```

Listing 4.2: My Caption

4.4.1.3 Comparing

We will compare JSON and XML features and performance with the study conducted by Saurabh and D’Souza [43]. There are apparent differences in the two data formats which affect the overall readability, extensibility, bandwidth performance, and ease of mapping. XML documents are easy to read, while JSON is obscure due to the parenthesis delimiters. XML allows for extended data types, while JSON is limited to a few data types. XML takes more bandwidth due to the metadata overhead, while JSON data is compact and use less amount of bandwidth.

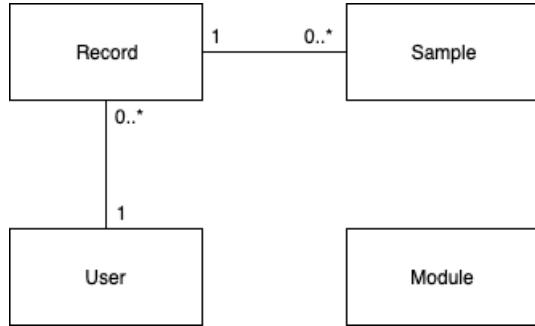


Figure 4.8: Modules

Moreover, a few benchmarks were conducted to measure memory footprint and parsing runtime when serializing and deserializing JSON and XML data. From the conclusion, in terms of memory footprint and parsing runtime, JSON performances better than XML but at the cost of readability and flexibility. While these format structures are applicable for transmitting data, choosing a format that is compact, human-readable, and a standard format that is extensible and scalable for future data is essential. In our design, we will be using the JSON format for transmission of the data.

4.4.2 Data Entities

Data entities are objects (e.g., things, persons, or places) that the system models and stores information about. Subsection about Storage introduces four data entities in our application (i.e., user, record, sample, and module). In Figure 4.8, the relation between the data entities are shown. Record and sample stores information about the recording, and are separated into two individual entities in order to reduce data redundancy and improve data integrity. Although, samples have a reference to its record so they can be associated with each other. A user stores biometrical information related to the user, the user in the application are patients. A record contains the state of the user's biometrical information at the time of the recording. In other words, the user's biometrical information can change over time (e.g., weight changes); therefore, capturing the exact biometrical information at the time of the recording is essential in the context of detecting sleeping illnesses with relation to the biometrical information. A module is independent of the other data entities and stores information about the name and the package name of the module-application. The package name is used to locate and launch the module-application.

In this Subsection, we will demonstrate the properties of each data entity in Nidra; storage structure of the entities, and illustration of the data structure for each entity.

id	name	description	monitorTime	rating	user	createdAt	updatedAt
1	Record #1	-	5963088	2.5	{...}	1554406256000	1554406256000

Table 4.1: Example entry in record table

4.4.2.1 Record

A record is a table in the database that stores metadata related to the recording session. In Table 4.1, an illustration of the structure of a record is shown, with an entry of dummy data. In Nidra, the fields in the table for a record describe the data that is stored, which is separated into:

- ID: Unique identification of a record, also a primary key for the entry.
- Name: A name of the record to easily recognize the recording.
- Description: A summary over the recording session provided by the user. It can be used to briefly describe how the recording session felt (e.g., any abnormalities during the sleep).
- MonitorTime: The recording session duration in milliseconds.
- Rating: Giving a rating on how the sleeping session felt, in a range between 0-5.
- User: User's biometrical information encoded into a JSON string format, in order to capture the state of the user at recording.
- CreatedAt: Date of creation of the recording in milliseconds (since January 1, 1970, 00:00:00 GMT).
- UpdatedAt: Date of update of the recording in milliseconds (since January 1, 1970, 00:00:00 GMT).

4.4.2.2 Sample

A sample is a single sensor reading containing data and metadata related to the recording session. Samples are stored separated from a record; however, they are linked with a field in the table. In Table 4.1, an illustration of the structure of a sample is shown, with an entry of dummy data. In Nidra, the fields in the table for a sample describe the data that is stored accordingly to the data provided by Flow SweetZpot (ref), which are separated into:

- ID: Unique identification of a sample, also a primary key for the entry.
- RecordID: An identification to its correlated record, also a foreign key.
- ExplicitTS: Timestamp of sample arrival based on the time in the sensor.
- ImplicitTS: Timestamp of sample arrival based on the time on the device.

id	recordId	explicitTS	implicitTS	sample
1	1	1554393086000	1554400286000	Time=0ms, deltaT=100, data=1906,1891,1884,1881,1876,1718,1690

Table 4.2: Example entry in sample table

id	name	packageName
1	OSA Predictor	com.package.osa_predicter

Table 4.3: Example entry in module table

- Sample: Sensor reading contains metadata and data according to Flow Sweetzpot. The sensor aggregates seven samples in a single sensor reading.

4.4.2.3 Module

A module is a table in the database that stores all modules installed by the user in the application. In Table 4.3, an illustration of the structure of a module is shown, with an entry of dummy data. In Nidra, the fields in the table contain the name of the module and the reference to the module and can be summarized as:

- ID: Unique identification of a module, also a primary key for the entry.
- Name: The name of the module-application.
- PackageName: The package name of the module-application, such that it can be launched from Nidra.

4.4.2.4 User

The user of the application is the patient, which provides biometrical data (e.g., weight, height, and age). The biometrical data is part of the application to enrich the record. The record captures the biometrical state of the user at the time of the recording in the context of detecting sleeping illnesses with the relation to the biometrical data. In Nidra, the user entity is not a part of the database, but as an object stored in the device. The design decision of this choice is because the application is limited to one user at the time, and it makes it convenient to capture the state of the object of the given time of recording. It is possible to create a table in the database for the user entity, and for each time the user changes the biometrical data insert it is a separate entry in the database. The record could then have a reference to the latest user entry. However, that increases the complexity of the system, and as of now, the design of the application is to store the user's biometrical data into an object on the device.

In Listing X, an illustration of the structure of a user is shown, with dummy data. In Nidra, the user object is structured in a JSON format and

```
1 {
2     "user": {
3         "name": "Ola Nordmann",
4         "age": 50,
5         "gender": "Male",
6         "height": 180,
7         "weight": 60
8     }
9 }
```

Listing 4.3: My Caption

contains biometrical data related to the user:

- Name: Name of the patient.
- Age: Age of the patient.
- Gender: Gender of the patient.
- Weight: Weight of the patient in kilograms.
- Height: Height of the patient in centimeters.

4.4.3 Data Packets

Data packets are parcels of data that Nidra receives from external applications (e.g., sensor wrappers) or send to other application (e.g., sharing). From the design choice In the Section above, the format of all the desired data should be according to the JSON format. In this Section.

4.4.3.1 Sharing

In Section 4.3.2, a proposal to the structure of exporting and importing data is discussed. Two of the components (Parse Formatted Content and Format File) uses JSON to either encode or decode the data. Listing 4.3 illustrate the content of the encoded data from our application to gain a broader understanding of how the data exchange in sharing operates. The attractive attributes from the encoding are the record (ref) and the samples. A record is an object that contains meta-data with name, number of samples, recording time, creation date, and user information. Samples are an array of objects that contains data, timestamp, and identification to correlated record.

```
1  {
2      "record": {
3          "id": 1,
4          "name": "Record 1",
5          "rating": 2.5,
6          "description": "",
7          "nrSamples": 6107,
8          "monitorTime": 5963088,
9          "createdAt": "Apr 4, 2019 9:30:56 PM",
10         "updatedAt": "Apr 4, 2019 9:30:56 PM",
11         "user": {
12             "age": 50,
13             "createdAt": "---",
14             "gender": "Male",
15             "height": 180,
16             "name": "Ola Nordmann",
17             "weight": 60
18         }
19     },
20     "samples": [
21         {
22             "explicitTS": "Apr 4, 2019 5:51:26 PM",
23             "implicitTS": "Apr 4, 2019 7:51:26 PM",
24             "recordId": 1,
25             "sample": "Time=0ms, deltaT=100, data
26                 =1906,1891,1884,1881,1876,1718,1690"
27         },
28         ...
29     ]
}
```

Listing 4.4: My Caption

4.4.3.2 Sensor Data

Sensor data is data acquisition through the Data Dispatching (section background). The data format discussed in (section background).

Chapter 5

Implementation

5.1 Android Components

The CESAR project is introduced in the Background Chapter; to summarize, the data dispatching module discovers and connects with sensor wrappers (each sensor source has its own sensor wrapper), and enables data acquisition to applications that subscribe to the data. As for this thesis, we operate with three different applications: Nidra, Data Stream Dispatching Module, and the sensor wrapper for Flow sensor kit.

Figure 5.1 illustrates the Android components (i.e., activity, service and broadcast receivers) for each applications. All of the applications run in a separate process on a device. In order to perform remote procedure calls (RPC) to application components that run remotely, we can use the IPC mechanism. In Android there are two mechanisms to enable IPC: (1) Binder enables a process to remotely invoke functions in another process; and (2) Intent a message passing interface allowing applications to send messages to each other.

5.1.1 Inter-Process Communication: AIDL

To perform IPC using Android Interface Definition Language (AIDL) [cite] we need to define a programming interface that both the client and the service agree upon. In order to communicate with processes, the data objects has to be decomposed into primitives that the operating system can understand, and marshall the objects across the boundary. The AIDL interface is defined in an .AIDL file, and located in the `src/` directory of the hosting service application (DSDM), and other applications that binds to the service (Nidra and sensor wrappers). It is important to have identical .AIDL files across the applications, otherwise the system will not recognize it as the same interface. In Listing 5.1, the interface is based on the functionality the hosting service application exposes (DSDM). In Nidra, some of the functionality is utilized to enable recording. More specifically,

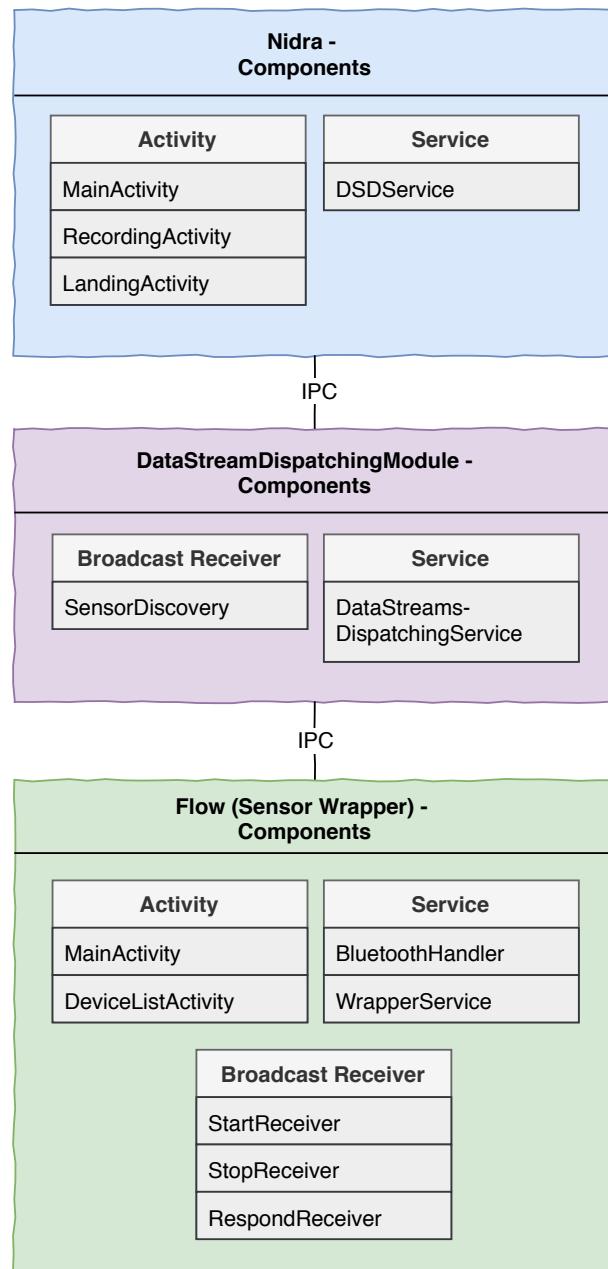


Figure 5.1: Applications components

`getPublishers()` method is used to get all of the sensors publishers (e.g., Bitalino provides multiple sensor capabilites ...), the `Subscribe(...)` and `Unsubscribe(...)` is used in order to subscribe and unsubscribe to a spesific sensor capability, and listing for events on the `putJson(...)`, which is used by the sensor wrappers to send data collected to DSDM, and further sent to all subscribing applications (i.e., Nidra).

```
1 // MainServiceConnection.aidl
2 package com.sensordroid;
3
4 interface MainServiceConnection {
5     void putJson(in String json);
6     int Subscribe(String capabilityId, int frequency, String
7                   componentPackageName, String componentClassName);
7     int Unsubscribe(String capabilityId, String
9                   componentClassName);
8     String Publish(String capabilityId, String type, String
9                    metric, String description);
9     void Unpublish(String capabilityId, String key);
10    List<String> getPublishers();
11 }
```

Listing 5.1: My Caption

5.2 Flow Sensor Kit Development

5.3 Implementation of Concerns

In Section (ref) we conceptualized the tasks, by decomposing the tasks into components and discussing various techniques and design decisions for implementation. In this Section, we will realize the discussion by implementing the tasks in Android.

5.3.1 Recording

Recording is the process of collecting and storing physiological signals from sensors over an extended period. To enable a recording, we need to establish a connection with the available sensors and store the samples retrieved by the sensors on the device. We will use the Data Stream Dispatching Module (hereafter: DSDM) which manages sensor discovery and sensor establishment to supported sensor sources. DSDM facilitates an interface for data acquistion, and the communication between the DSDM and Nidra occurs over IPC. Moreover, we will implement functionality to check for connectivity during record, to ensure the sensors are collecting

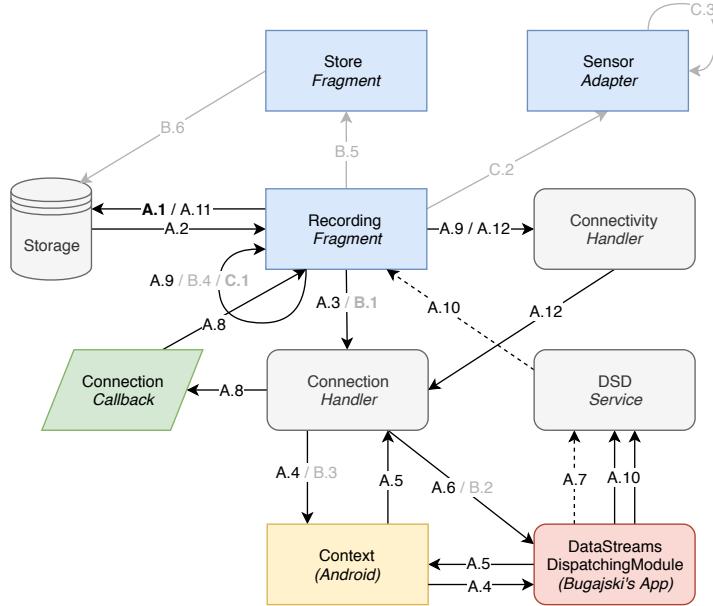


Figure 5.2: Implementation of recording functionality: (A) Start a Recording

data to at appropriate rate. At the end of recording, we will store metadata related to the recording session and finalize the recording process.

The functionality of recording can be divided into three actions: (A) start a recording; (B) stop a recording; and (C) display recording analytics. In the following Subsection for recording, we will review the steps to enable these actions.

5.3.1.1 Action A: Start a Recording

In Figure 5.2, an illustration of the component interactions are shown. Action A is to start a recording by connecting and starting data aquisition with the use of DSDM, and to ensure persistent connectivity with the sensor sources. The steps and interactions for this action are:

- A.1 Commence the recording by creating a record entity and inserting it into the database. An empty record has to be inserted into the database in order to associate new samples to the recording.
- A.2 Once the record is inserted into the storage, unique identification is returned.
- A.3 **ConnectionHandler** manages the establishment, connection, and disconnection of the IPC between Nidra and DSDM. It starts by establishing a connection to the DSDM:

```

1     Intent intent = new Intent(MainServiceConnection.class.
2         getName());
3     intent.setAction("com.sensordroid.ADD_DRIVER");
4     intent.setPackage("com.sensordroid");
5     context.bindService(intent, serviceCon, Service.
6         BIND_AUTO_CREATE);

```

Listing 5.2: My Caption

MainServiceConnection is the AIDL file (discussed in IPC).

- A.4 We bind to service by using the BindService. If the service is offline, the flag Service.BIND_AUTO_CREATE will ensure the service is started. BindService allows components to send requests, receive responses, and perform inter-process communication (IPC) [CITE].
- A.5 Once the service has bound, we can proceed to communicate with the DSDM.
- A.6 The ConnectionHandler proceeds to initialize the connection with the sensor through the DSDM. A request to the DSDM for available publishers with getPublishers() is made, to retrieve all available sensor publishers connected to the DSDM. Occasionally, the DSDM uses extended time to discover all of the active sensors connected to the device; therefore, we have an interval that checks whether DSDM has any available sensors connected.
- A.7 Moving on, a request to the DSDM to Subscribe to a sensor is made. In the Subscribe function, a reference to the package name and a service object (DSDService) from Nidra is sent. The service object is where all of the parcels of data from the DSDM is received.
- A.8 A callback to RecordingFragment with the available sensor publishers is made.
- A.9 The recording has now started, and a timer to measure the time spent on the record is started.
- A.10 The ConnectivityHandler is the component we discussed in Section ... , which checks for sample arrival and activity to reconnect with the sensor. The ConnectivityHandler is implemented with a Handler with a PostDelay. If the event for the PostDelay is triggered, it is equivalent for a sample not being acquired from the sensor.
- A.11 Periodically, the DSDM receives samples from the subscribed sensor. DSDM forwards the sample from the sensor to the service object (DSDService) on the \verbputJson! function. The DSDService uses a LocalBroadcastManger to send bundles of data across components in the application. The RecordingFragment is listening for the event, and receives all of the data.
- A.12 The data received from the sensor (through the DSDM, to DSDService, and LocalBroadcast), is stored as a new sample entry with the

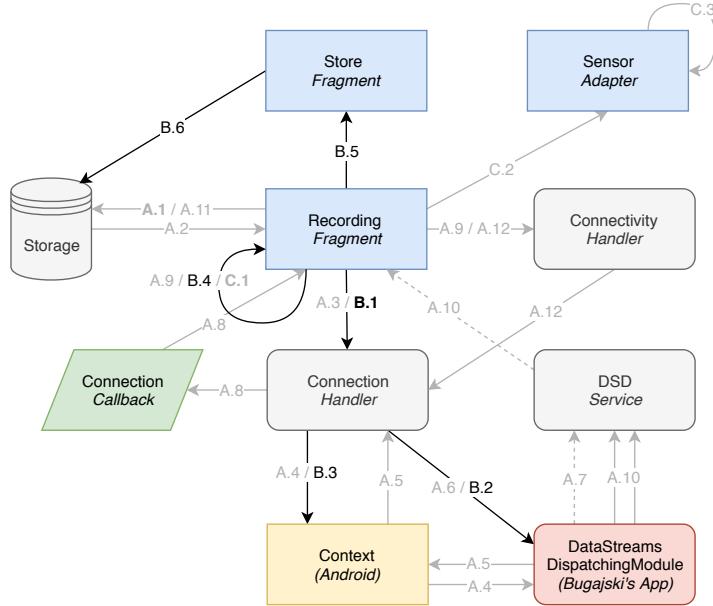


Figure 5.3: Implementation of recording functionality: (B) Stop a Recording

current recording identification associated with it.

5.3.1.2 Action B: Stop a Recording

In Figure 5.3, an illustration of the components to stop a recording is shown. The action is based on the user input to stop the recording session. For the user, the recording has terminated and the user is presented with a screen to provide information regarding the recording (e.g., title, description and rating). For the application, has to notify and unsubscribe to the connected sensors through DSDM, as well as disconnect from the DSDM service. In addition, manage the metadata the user provides at the end of the recording session, and storing recording metadata on the device. The steps and interaction for this action are:

- B.1 The user of application decides when to stop a recording with a press of a button. The event to stop recording is sent to the `ConnectionHandler`.
- B.2 A call to the `Unsubscribe` function on the subscribed subscribe with the service object is sent to the `DataStreamsDispatchingModule` (hereafter: DSDM). The DSDM has to ensure to unsubscribe unpublis the sensor in order to signal the sensor to stop sampling.
- B.3 The IPC connection between Nidra and DSDM is discontinued by unbinding the service.
- B.4 The estimated time of recording is calculated, and a transition from `RecordingFragment` to `StoreFragment` is made to finalize the recording with extra information (e.g., title, description, and rating).

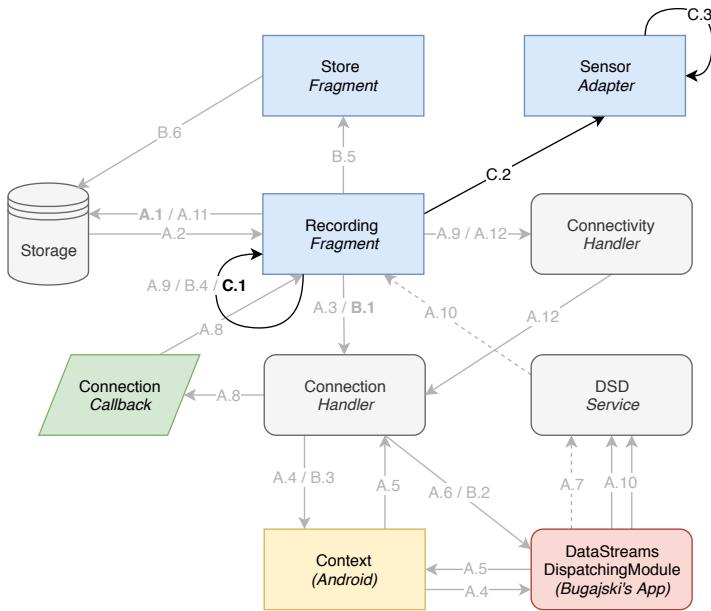


Figure 5.4: Implementation of recording functionality (C)

- B.5 The `StoreFragment` uses the record identification retrieved on recording (A.1) in order to enrich the record with statistics and user-defined metadata. The statistics are the monitoring time, number of samples during recording, and retrieving the current state user biometrical data and storing it in the record. The user-defined metadata are the title of the recording, a description enabling the user to add a note to the recording, and a rating between 1–5 (to give a rating on how the recording felt).
- B.6 The modified record is updated in the database, and the user is transitioned to the `FeedFragment`.

5.3.1.3 Action C: Display Recording Analytics

During a recording, the user can view the analytics for the recording. More specifically, the user can see the available connected sensors, and a graph of respiration (based on the samples acquisitions from the sensor) in real-time. In Figure 5.4, an illustration of the components display recording analytics is shown, and the steps and interaction for this action are:

- C.1 The data is graphically represented as an intractable time-series graph. By using the Graph library, we can in similarities to the implementation of the analytics concern (see Subsection), implement a graph to illustrate the respiration data to the user.
- C.2 In addition to the time-series graph, we have a list of sensors that are sampling in the recording.

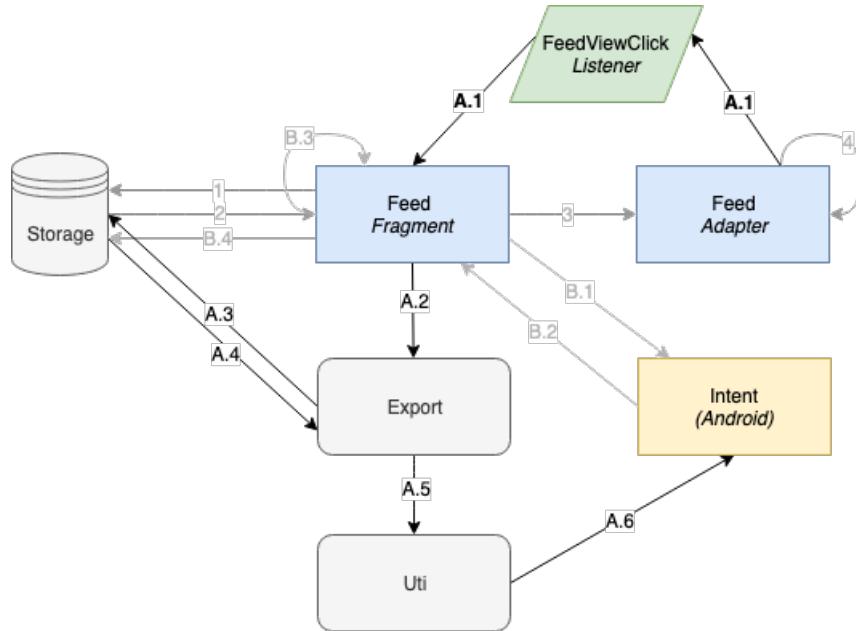


Figure 5.5: Implementation of sharing functionality (A): Exporting one or all Records

C.3 The Sensor Adapter populates and illustrates the connected sensor to the user.

5.3.2 Sharing

Sharing enables users to transmit records across application over a media. The functionality of sharing is separated into two concerns, namely importing and exporting records. Hence, the actions for sharing are separated into: (A) exporting one or all records; and (B) import a record from the device.

Before a user can take these actions, the records from the database have to be presented. The Feed Fragment contains a RecyclerView which populates the records into inside the Feed Adapter (steps: 1-4). The adapter contains all the interactions and the event handling (i.e., button event listener for exporting) for a single record.

In this Subsection, we will take a look into the steps that are taken to enable the actions:

5.3.2.1 Action A: Exporting one or all Records

In Figure 5.5, an illustration of the steps to export one single recording is shown. However, the Feed Fragment has an option to export all records; therefore, by disregarding the first step (A.1), the same structure applies to

export all records. In essence, exporting consists of bundling the records and corresponding samples into a formated file, and prompting the user with options to select a media (e.g., mail) for transmission. The steps can be narrowed down to:

- A.1 Upon an event for exporting a selected record in Feed Adapter, the record information is sent to the Feed Fragment through the callback reference (`onRecordAnalyticsClick`) between these components. The record information will be used to determine the corresponding samples for the record.
- A.2 The Feed Fragment delegates record information to the `export` method inside of the Export class. The class is responsible for enabling exportation.
- A.3 An operation to retrieve all samples related to the record with the use of the `SampleViewModel` is done.
- A.4 The `export` method retrieves all of the samples related to the record. Next, the record and the samples are encoded into an exportable JSON format (Ref: Data Format). To enable the sharing interface provided by Android, the content has to be stored on the device. Thus, the encoded data is written into a file on the device, with a filename of `record_(current_date).json`, and the next component uses the reference to the file location.
- A.5 The encoded file is retrieved with the use of `FileProvider` (facilitates secure sharing of files [ref]). The code for this step are

```
1 static void shareFileIntent(Activity a, File file) {  
2  
3     Uri fileUri = FileProvider.getUriForFile(a.  
4         getApplicationContext(), a.getApplicationContext().  
5         getPackageName() + ".provider", file);  
6  
7     Intent iShareFile = new Intent(Intent.ACTION_SEND);  
8     iShareFile.setType("text/*");  
9     iShareFile.putExtra(  
10        Intent.EXTRA_SUBJECT, "Share Records");  
11     iShareFile.putExtra(Intent.EXTRA_STREAM, fileUri);  
12     ...  
13     a.startActivity(  
14         Intent.createChooser(iShareFile, "Share Via"));  
15 }
```

Listing 5.3: My Caption

- A.6 The user is displayed with a popup interface with several options to share the file over a media. An illustration of the layout can be found

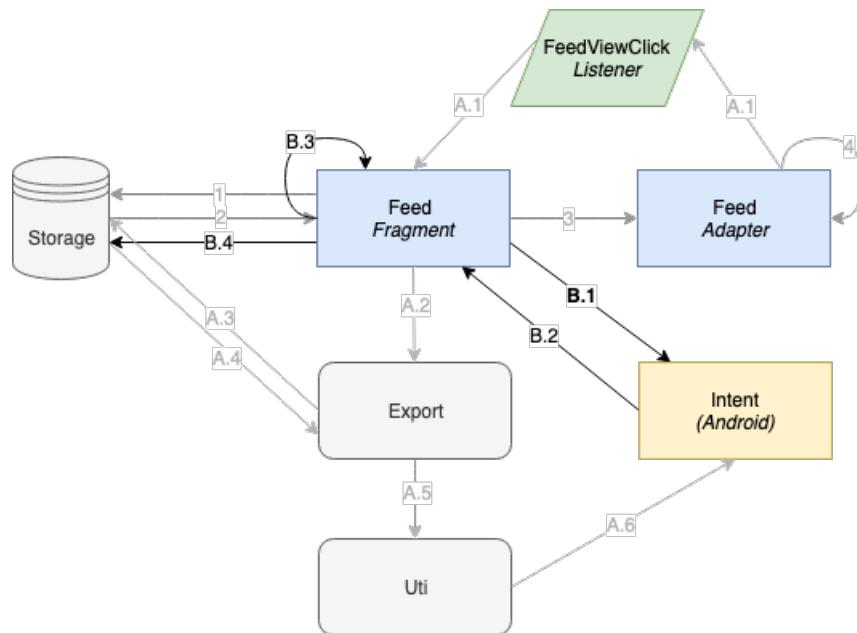


Figure 5.6: Implementation of sharing functionality (B)

in Section Representation.

5.3.2.2 Action B: Import a Record from the Device

In Figure 5.6, an illustration of importing a record from the device is shown. Importing consists of locating the formatted file (the user has to obtain the file and store it on the device beforehand), parsing the content in the file, and storing the data respective to the users database. The steps can be narrowed down to:

- B.1 The user requests to view the import record interface. The interface is provided by Android, and allows the user to select particular kind of data on the device (ref). The code for this action is:

```

1 private void importRecords() {
2     Intent intent = new Intent(Intent.ACTION_GET_CONTENT);
3     intent.setType("/*/*");
4     startActivityForResult(intent, 1);
5 }

```

Listing 5.4: My Caption

- B.2 Once the user has selected the desired file, the method `onActivityResult` inside of `Feed Fragment` is called, and location of the selected file can be located.
- B.3 The file location is an obscured path to the file on the device; thus,

parsing the path with the use of `Cursor` method has to be done. After the absolute path is found, the data is decoded accordingly to the data format, and the records are sent back to Feed Fragment.

- B.4 The necessary record information and the samples are extracted from the decoded data, and are inserted into the users database.

5.3.3 Modules

Modules are standalone application, that provides data enrichment and extended functionality to the application. The modules leverages the records and samples to analyze, evaluate or detect sleeping disorders. In order to add and launch a module in Nidra, we need the modules package name. The package name and the name of the module-application can be obtained in Android. Thus, the actions to enable modules in the application are: (A) add a module; and (B) launch a module.

Before a user can take these actions, the records from the database have to be presented. The Feed Fragment contains a `RecyclerView` which populates the records into inside the Feed Adapter (steps: 1-4). The adapter contains all the interactions and the event handling (i.e., button event listener for exporting) for a single record.

In this Subsection, we will take a look into the steps that are taken to enable the actions.

5.3.3.1 Action A: Add a Module

In order to add a new module, the user has to install the module-application on the device on beforehand. By listing through the installed application on the device, the user can select the desired module to be added in Nidra. In Figure 5.7, an illustration of adding a module is shown, and the steps can be narrowed down to:

- A.1 Upon an event for adding a new module in Modules Adapter, the Feed Fragment is notified through the callback reference (`onNewModuleClick`) between these components.
- A.2 The Modules Fragment launches a custom Android dialog, which will list all of the installed applications on the device.
- A.3 The Apps Adapter will fetch all of the applications that are not a system package, already installed module, or the current application (Nidra). Next, the the adapter for the dialog will be populated with the eligible applications.
- A.4 Once the user has selected the desired module-application, an event to the Modules Fragment through the callback reference `onAppItemClick` between these components are made. The callback

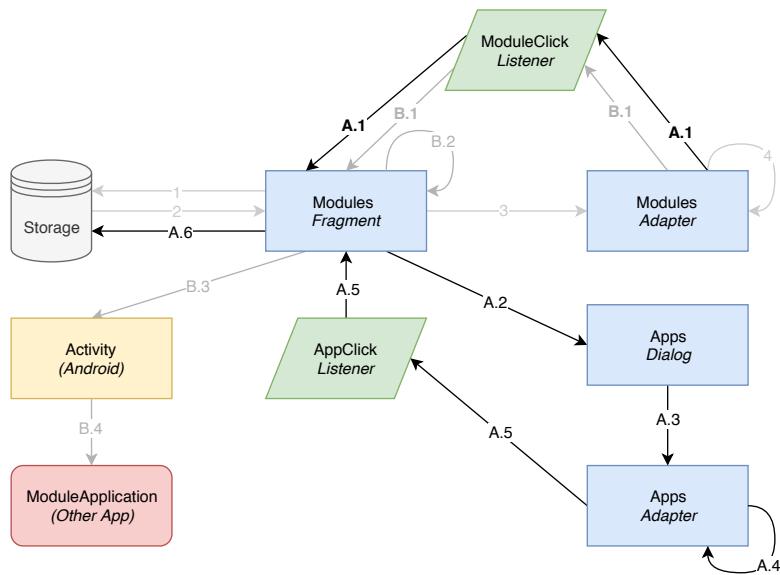


Figure 5.7: Implementation of module functionality(A): Add a Module

contains an object with the `PackageInfo` for the selected module-application.

- A.5 The dialog is dismissed, and the application name and package-name are extracted from the `PackageInfo` for the selected module-application.
- A.6 Furthermore, the acquired information is stored in our database for modules through the DAO interface.

5.3.3.2 Action B: Launch a Module

A module is launched in a separate process, due to Android prohibits launching for other applications inside of an application. All added modules are listed and presented to the user in a separate screen, and on launch of a module, all of the data that Nidra obtains from recordings, are encoded into a JSON format and bundled with the launch of the module. In Figure 5.8, an illustration of launching a module, and the steps can be narrowed down to:

- B.1 Upon an event for launching a module in `Module Adapter`, the packagename of the module is sent to the `Modules Fragment` through the callback reference (`onLaunchModuleClick`) between these components. The packagename will be used to launch the module-application.
- B.2 All of the records and samples on the device for the user, is bundled and formatted into a JSON, and launched:

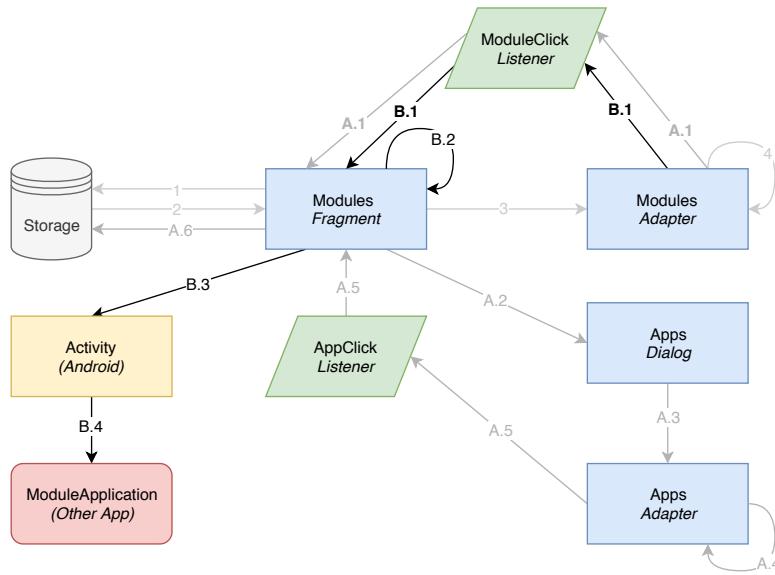


Figure 5.8: Implementation of module functionality(B): Launch a Module

```

1 public void onLaunchModuleClick(String packageName) {
2     Intent moduleApplication = context.getPackageManager() .
3         getLaunchIntentForPackage(packageName);
4
5     if (moduleApplication == null) return;
6
7     String data = formatAllRecordsToJSON();
8
9     Bundle bundle = new Bundle();
10    bundle.putString("data", data);
11
12    moduleApplication.putExtras(bundle);
13
14    startActivity(moduleApplication);
}

```

Listing 5.5: My Caption

B.3 The activity uses the data provided in the Intent that includes the packagename (the name of the module-application to determine the correct application).

B.4 The selected module is then launched, and presented to the user. The user can at anytime press the back button, to return to Nidra.

5.3.4 Analytics

Analytics is the part of illustrating and analyzing the records. In Nidra, the analytics part of the implementation is limited to a time-series graph

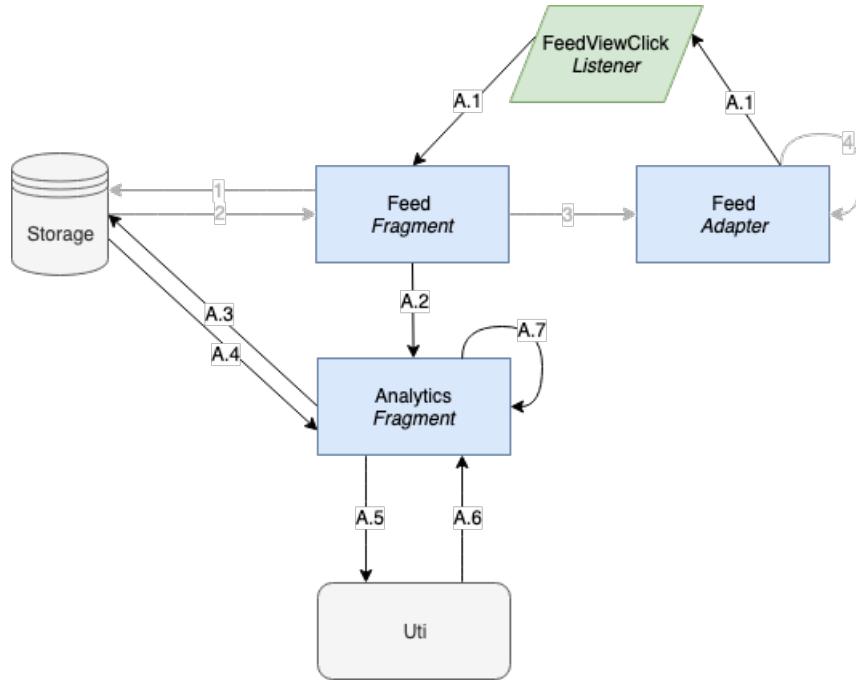


Figure 5.9: Implementation of analytics functionality (A): Display a Graph for a Single Record

for a single record. However, there are possibilities of extending the Analytics Fragment with other graphs based on the current structure. The current action for analytics is A) display a graph for a single record to the user.

Similar to sharing, the records from the database have to be presented. The Feed Fragment contains a RecyclerView which populates the records into inside the Feed Adapter (steps: 1-4). The adapter contains all the interactions and the event handling (i.e., button event listener for analytics) for a single record.

In this Subsection, we will take a look into the steps that are taken to enable the action.

5.3.4.1 Action A: Display a Graph for a Single Record

Nidra provides a simple time-series graph of respiration data obtained during the recording. The graph data is plotted into a Library, which enables interactions (e.g., zoom and scrolling) on the data. The X-axis is the respiration value based on the Y-axis time of sampling. In Figure 5.9, an illustration of displaying a graph shown. The steps can be narrowed down to:

- A.1 Upon an event for analytics on a selected record in Feed Adapter, the record information is sent to the Feed Fragment through the callback

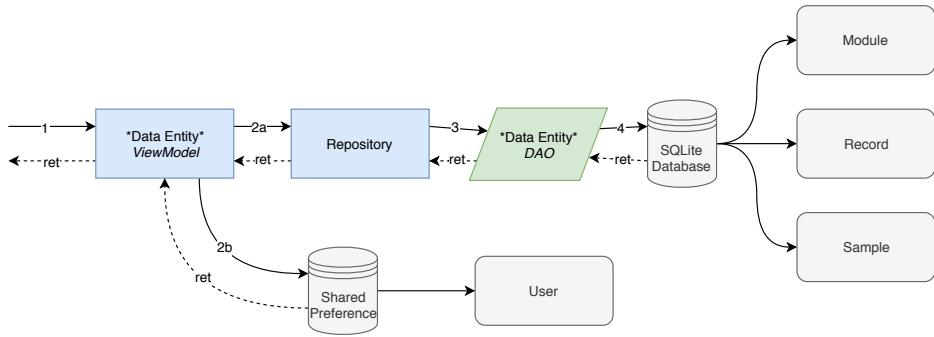


Figure 5.10: Entity Relationship Diagram

reference (`onRecordAnalyticsClick`) between these components. The record information will be used to determine the corresponding samples for the record.

- A.2 A new instance of the Analytics Fragment is created, and an transition from the Feed Fragment to the Analytics Fragment is made. Alongside, the record information is transmitted.
- A.3 An operation to retrieve all samples related to the record with the use of the `SampleViewModel` is done.
- A.4 The Analytics Fragment retrieves all of the samples related to the record. The samples has to be structured according to the graph library to display an interactive time-series graph. (GraphView (ref)).
- A.5 Each sample has to be extracted from the sample-data, acccordingly to the sensor data structure.
- A.6 The sample value is returned, and inserted into an array over datapoints used in the graph.
- A.7 The use is presented with a graph, which is interactable. The Y-axis has the sample value on given time (in HH:MM:SS) on the X-axis. The graph library enables interactions (e.g., zooming and scrolling) the user can do, to gain a better understanding of the recording.

5.3.5 Storage

Storage facilities persistent data which remain available after application termination. In Nidra, there are four individual data entities (i.e., record, sample, module, and user). In the Subsection about Data Entities, we discussed the design choices of the individual data entity. The standard CRUD operations on each data entities are: insert, update, and delete. Also, the user has an operation to retrieve the biometrical data, module and record have an operation to retrieve all of the entries in the database, and samples have an operation to retrieve all of the entries corresponding to a record.

Android Room provides an abstract layer over SQLite to enable easy database access [Cite]. In Figure 5.10, the flow for accessing and retrieving the data from the database based on the Android Room architecture is shown and can be described as:

- 1 Each data entity has a `ViewModel` where all of the CRUD operation (e.g., insert, update, delete, or retrieve) goes through. A view model is designed to store and manage UI-related data in a conscious way, that allows data to be persistent through configuration changes (e.g., screen rotations) [CITE].
- 2a The predefined operations point to the repository. Repository modules handle data operations and provide an API which makes data access easy. A repository is a mediator between different data sources (e.g., database, web services, and cache) [CITE]. In Nidra, the only data source is the database, but repository facilities future data sources.
- 2b The storage of the user is not in the database; however, in a `SharedPreference` on the device. Shared preference points to a file containing key-value pairs and provides methods to read and write. The location of the user's shared preference is `no.uio.cesar.user_storage`.
- 3 Each data entity (disregarding user) has a data access object (DAO), where the SQL operations to the database are defined.
- 4 Based on the operation, the data is accessed and retrieved from the SQLite database.

5.3.6 Presentation

In this Subsection we will present the user-interface (UI) based on the functionality of recording, sharing, module and analytics. The interface is developed based on the design decisions of the application, as well as creating a user-interface that is simple and efficient for a user to interact with. We try to limit the actions the user can take on a screen, to make the application simpler to understand and comprehend.

5.3.6.1 UI: Recording

Figure 5.11 shows the screen for: (A) during a recording, (B) the real-time analytics, and (C) finalizing the recording.

- A The screen has a ripple effect to indicate the state of the recording to the user. There are two types of ripple colours; a blue ripple effect for samples acquisition, and a grey ripple effect if the sensor has disconnected. The ripple effect is only active if the screen is turned on, in order to preserve battery life. During disconnects between the

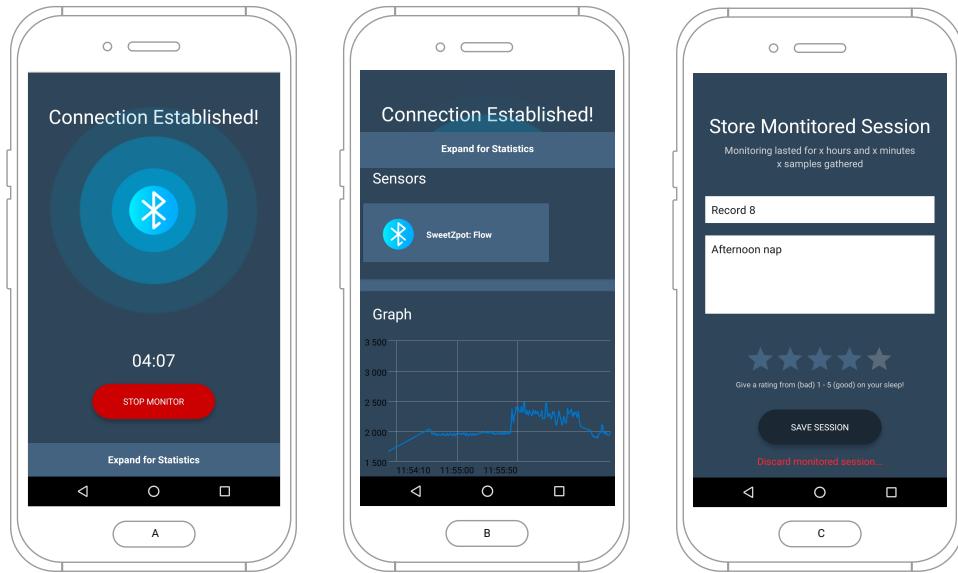


Figure 5.11: The recording screen displayed to the user; with the screen: (A) during a recording, (B) real-time analytics, and (C) finalizing the recording

sensor and the device, the user provide no extra input to resolve the issue. However, if the numbers of attempts is reached, the user is presented with the finalizing screen (C). Also, this screen display the elapsed recording time, and has a visible button to stop the recording.

- B The user can extend the interface for viewing statistics regarding the recording session. Currently, the interface lists all of the available sensor sources, also a real-time interactable graph of respiration.
- C The finalizing screen allows users to specify the title and description of recording, to make it distinguishable, in addition to giving a feedback on recording so the doctors/researchers can review it. Also, the user can give a rating between 1–5 (where 1 is bad, and 5 is good) to rate the sleeping.

5.3.6.2 UI: Sharing

Figure 5.12 shows the screen for: (A) option to import or export, (B) the media selection for exporting, and (C) the file selection for importing.

- A The feed screen is where all of the records are presented to the user. The user can choose to export one single record, or export all.
- B By pressing export all, an overlay with a Android provided sharing screen is presented, where the user can choose a media to export the file on.
- C By pressing import, an overlay with downloaded files on the users device is presented. The user can press on the desired file, and the

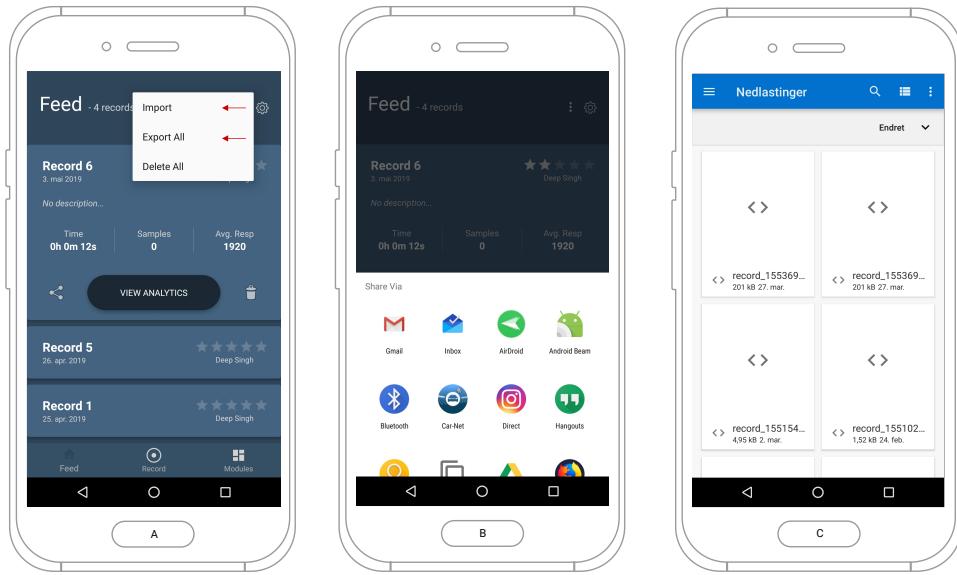


Figure 5.12: The sharing screen displayed to the user

file will be parsed and added to the users collection of records.

5.3.6.3 UI: Modules

Figure 5.13 shows the screen for: (A) list without any modules, (B) list of installed application, and (C) list with modules.

- A The modules screen is shows all of the added modules to the users.
- B The user can press the *add new module* button, in order to be presented with a list of all installed applications.
- C Once the user has selected a module-application to be added in Nidra, the list of modules are updated and presented to the user.

5.3.6.4 UI: Analytics

Figure 5.14 shows the screen for: (A) single record in the list, and (B) the analytics for the record

- A The user can expend records in order to view more functionality and statistics of the record. One of the functionalities are to view the analytics of the record.
- B An interactable graph is presented to the users, and the graph is populated with the samples obtained from the recording for selected records. The graph shows the respiration value (Y-axis) on given time (X-axis) of sampling.

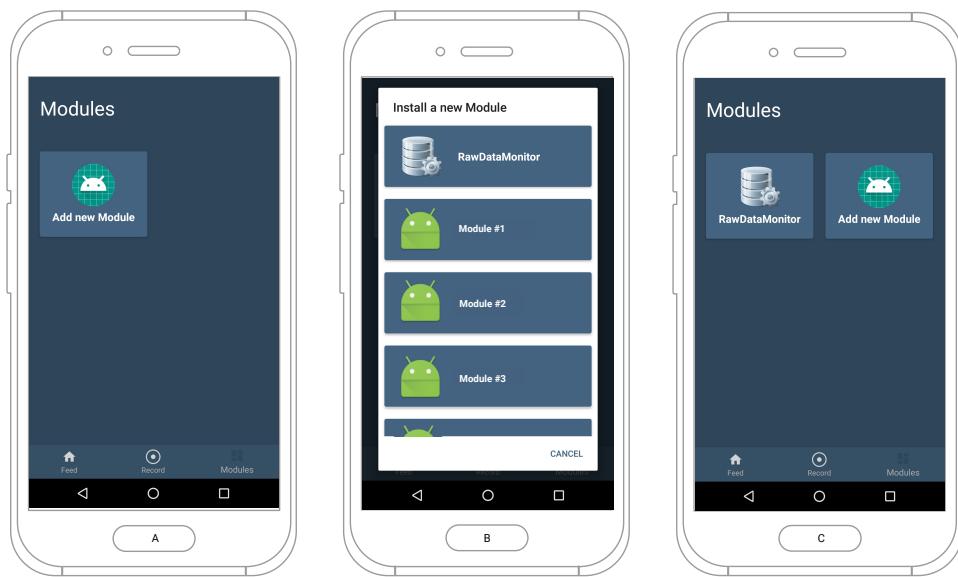


Figure 5.13: The module screen displayed to the user

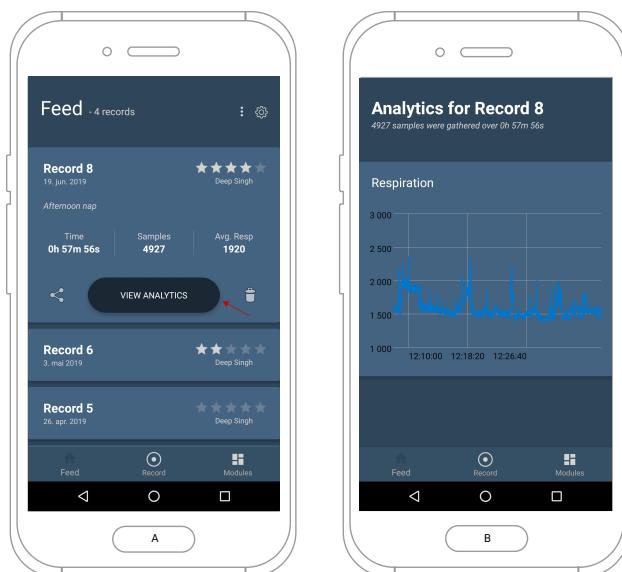


Figure 5.14: The analytics screen displayed to the user

Part III

Evaluation and Conclusion

Chapter 6

Evaluation

6.1 Experiments and Measurements

6.1.1 Experiment A: Orchestral Concert to Analyze Musical Absorption using Nidra to Collect Breathing Data

The experiment was conducted in collaboration with master student Joachim Dalgard at the University of Oslo at the faculty of *RITMO: Centre for Interdisciplinary Studies in Rhythm, Time and Motion*.

The goal of the experiment was to analyze musical absorption, which is a state an individuals ability and willingness allows music to draw them into an emotional experience and becomes unaware of time and space. In order to analyze the effect of musical absorption on individuals, we gathered 20 participants who were experienced listeners with musical education. The participants attended an orchestral concert by Richard Strauss' Alpine Symphony—a symphonic poem that portrays the experience of eleven hours spent climbing an Alpine mountain—that lasted around 50 minutes at Oslo Concert Hall on third of April and fourth of April 2019.

The participants were divided into two groups to attend the concert on the two dates. Each participant was equipped with a wireless electromyographic sensor from DELSYS in order to measure heart rate, and a Flow sensor kit to measure respiration during the concert. RITMO had multiple Flow sensors for disposal; however, they had no suitable mobile application that could record with these sensors. Also, with their equipment they experienced that Flow sensor kits had a tendency to disconnect every 10-15 minutes, resulting in fragmented recordings for a single session. Therefore, they reached out to *Institute for Informatics* in hopes of a solution. Our application was a suiting match for both parties, hence a collaboration was formed. We arranged for six Android devices and reached out to the participants to bring their Android devices if they had one. There were ten participants in each group and six assessed

Model	Samsung Galaxy S9	OnePlus 3T	Google Pixel XL
Operating System	Android 8.0	-	Android 9.0 & Android 7.1.2
Chipset	Exynos 9810	Qualcomm MSM8996 Snapdragon 821	Qualcomm MSM8996 Snapdragon 821
CPU	Octa-core	Quad-core	Quad-core
GPU	Mali-G72 MP18	Adreno 530	Adreno 530
RAM	4 GB	6 GB	4 GB
Battery	Li-Ion 3000 mAh	Li-Ion 3400 mAh	Li-Ion 3450 mAh
Bluetooth	5.0, A2DP, LE, aptX	4.2, A2DP, aptX HD, LE	4.2, A2DP, LE, aptX

Table 6.1: Device models used during the concert

devices. As a precaution, we decided to give out the devices to the participants who scored highest on a test performed on beforehand.

The motivation for this experiment in regards to Nidra was to test the application in a real-life and crowded environment. During the concert, there were approximately 800 attendees on the first day and approximately 1500 attendees on the second day. We can assume that most of the attendees had a mobile device, and a few shares of them had BlueTooth activated on the device. Based on these estimates, we were able to replicate an environment (on a larger scale) where other devices might interfere with the signals between the collecting sensor and the application. Also, we were able to install the application on multiple mobile devices with different Android OS versions and put the application in the hands of the participants.

6.1.1.1 Preperations

In order to partake in the experiments, we had to ensure that the mobile devices were configured correctly. Also, stress testing the application to prevent ...

Device Configuration The device models in our disposal had to be configured with the applications to enable recording on Nidra. First, the data stream dispatching module had to be installed on the devices. Second, the sensor wrapper for the Flow sensor kit had to be configured with the respective sensor (in order to reduce the time on setup the sensor and application to the participant). Lastly, the Nidra application had to be initiated with appropriately user information, in order to distinguish each participants with the sensor and device model. In Figure 6.1, the device models are listed with their spesifications, and the OS ...

Obstacle One of the pressing concerns at the time of the experiment was to ensure a persistent connection between the sensor and application—uninterrupted by sensor disconnections and disruptions—in order to collect meaningful data. At the time of development, there was no mechanism to reconnect with the sensors on disconnect or disruption during recording. Therefore, the mechanism had to be developed, otherwise, the recording would not be able to capture the respiration of the participant through the whole concert.

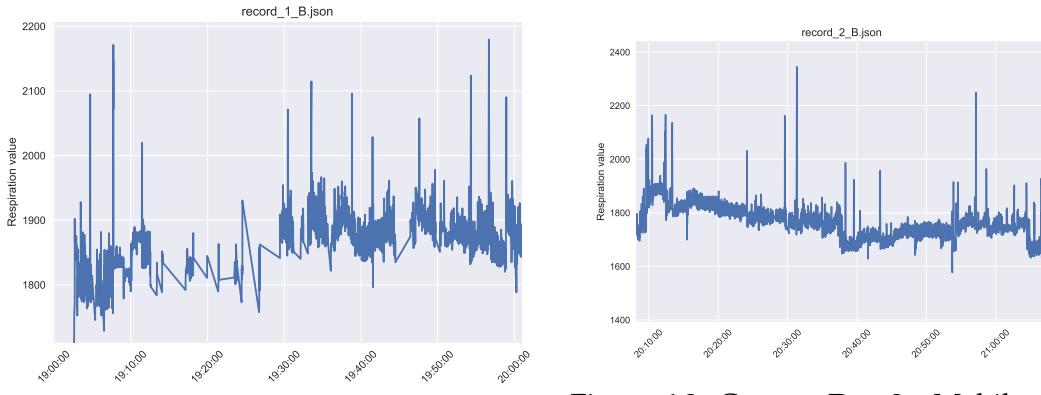


Figure 6.1: Concert Day 1—Mobile Device B

Figure 6.2: Concert Day 2—Mobile Device B

Body Placement The respiration value is based on the participant body circumference and changes are associated with breathing. The participant was instructed to place the sensor around their thorax (just below the armpits), in order to measure the expansion and contraction of the rib cage, to measure the respiration flow. Also, the respiration can vary through the recording based on body position (e.g., sitting relax or tense in a chair).

6.1.1.2 Results

We had thirteen mobile devices combined for the both dates—the extra device was from one of the participants—however, the application crashed on one of the mobile devices during the recording, rendering in an incomplete recording on one of the mobile devices. Therefore, we have access to twelve recordings from the concerts. Figure 6.1 and Figure 6.2 present two of these twelve recordings (rest can be found in Appendix B) in a time-series graph that are of most interest to us. The Y-axis presents the respiration (breathing) value, and the X-axis the time of respiration value acquisition—day 1 of the concert started at time of 19:00 and ended at 20:00, while day 2 started at time of 20:10 and ended 21:00. The graphs are plotted with Python and the library Matplotlib in order to analyse and evaluate the graphs; however, an identical graph is also presented in the application.

6.1.1.3 Analysis & Discussion

To begin with, we will investigate all of recordings by analyzing the respiration values. We are specially interested in finding occurrences of disconnects—when the samples are stagnant in longer period, in contrasts to spikes which emulates breathing by contraction and extraction of the body—and the time it takes the application to reconnect with the sensor.

Model	Samples	Loss Count	Loss Percentage	Times Disconnected	Time Disconnected
A	5145	0	0 %	0	00:00
B	3363	1782	34 %	13	19m:14s
C	4189	956	19 %	7	8m:20s
D	3501	1644	32 %	5	18m:30s
E	5144	1	0.02 %	-	-
F	5145	0	0 %	0	00:00

Table 6.2: Day 1—Duration: 1 hour & Roof Samples: 5145

Model	Samples	Loss Count	Loss Percentage	Times Disconnected	Time Disconnected
A	4286	2	0.05 %	-	-
B	4161	127	3 %	4	1m:25s
C	4286	2	0.05 %	-	-
D	2576	1712	40 %	7	24m:35s
E	4285	3	0.06 %	-	-
F	4288	0	0 %	0	00:00

Table 6.3: Day 2—Duration: 50 mins. & Roof Samples: 4288

To conclude this experiment,

6.1.2 Experiment B: 8-hours recording

6.1.3 Experiment C: User-Experience

6.1.4 Experiment D: Creating a Simple Module

6.2 Main Findings

Chapter 7

Conclusion

Appendix

Appendix A

Power Data

Bibliography

- [1] Ecma International 2017. *The JSON Data Interchange Syntax*. URL: <http://www.ecma-international.org/publications/files/ECMA-ST/ECMA-404.pdf> (visited on 01/06/2019).
- [2] *Activities*. Android Developer. URL: <https://stuff.mit.edu/afs/sipb/project/android/docs/guide/components/activities.html> (visited on 23/06/2019).
- [3] S. Alqassim et al. ‘Sleep Apnea Monitoring using mobile phones’. In: *2012 IEEE 14th International Conference on e-Health Networking, Applications and Services (Healthcom)*. Oct. 2012, pp. 443–446. DOI: 10.1109/HealthCom.2012.6379457. (Visited on 26/06/2019).
- [4] Android. *Secure an Android Device*. URL: <https://source.android.com/security> (visited on 30/05/2019).
- [5] *Android Interface Definition Language (AIDL)*. Android Developer. URL: <https://developer.android.com/guide/components/aidl> (visited on 24/06/2019).
- [6] *Application Fundamentals*. Android Developer. URL: <https://stuff.mit.edu/afs/sipb/project/android/docs/guide/components/fundamentals.html> (visited on 23/06/2019).
- [7] *Binder*. Android Developer. URL: <https://developer.android.com/reference/android/os/Binder> (visited on 24/06/2019).
- [8] *Bluetooth low energy overview*. Android Developer. URL: <https://developer.android.com/guide/topics/connectivity/bluetooth-le> (visited on 25/06/2019).
- [9] Daniel Bugajski. ‘Extensible data streams dispatching tool for Android’. In: (201).
- [10] CESAR: *Using Complex Event Processing for Low-threshold and Non-intrusive Sleep Apnea Monitoring at Home*. UiO: Department of Informatics. URL: <https://www.mn.uio.no/ifi/english/research/projects/cesar> (visited on 21/06/2019).
- [11] World Wide Web Consortium. *Extensible Markup Language (XML) 1.0 (Fifth Edition)*. URL: <https://www.w3.org/TR/REC-xml/> (visited on 01/06/2019).

- [12] *Content Providers*. Android Developer. URL: <https://stuff.mit.edu/afs/sipb/project/android/docs/guide/topics/providers/content-providers.html> (visited on 23/06/2019).
- [13] Abe Crystal and Beth Ellington. 'Task analysis and human-computer interaction: approaches, techniques, and levels of analysis'. In: *AMCIS*. 2004, pp. 2–3. URL: <https://pdfs.semanticscholar.org/fbd2/b61c998cb3ad8427759a45370a02d9338c31.pdf> (visited on 02/05/2019).
- [14] *Dark theme*. Material Design. URL: <https://material.io/design/color/dark-theme.html#usage> (visited on 04/06/2019).
- [15] *Data and file storage overview*. Android Developer. URL: <https://developer.android.com/guide/topics/data/data-storage> (visited on 25/06/2019).
- [16] *Data and file storage overview*. Android Developers. URL: <https://developer.android.com/guide/topics/data/data-storage> (visited on 05/06/2019).
- [17] *Fact check: Is smartphone battery capacity growing or staying the same?* Android Authority. URL: <https://www.androidauthority.com/smartphone-battery-capacity-887305/> (visited on 30/05/2019).
- [18] *FileProvider*. Android Developer. URL: <https://developer.android.com/reference/android/support/v4/content/FileProvider> (visited on 23/06/2019).
- [19] *Fragments*. Android Developer. URL: <https://developer.android.com/guide/components/fragments> (visited on 23/06/2019).
- [20] R. Fu, Z. Zhang and L. Li. 'Using LSTM and GRU neural network methods for traffic flow prediction'. In: *2016 31st Youth Academic Annual Conference of Chinese Association of Automation (YAC)*. Nov. 2016, pp. 324–328. DOI: 10.1109/YAC.2016.7804912. (Visited on 07/06/2019).
- [21] H Gray Funkhouser. 'Historical development of the graphical representation of statistical data'. In: *Osiris* 3 (1937), pp. 269–404. URL: <https://www.journals.uchicago.edu/doi/pdfplus/10.1086/368480> (visited on 07/06/2019).
- [22] Svein Petter Gjøby. 'Extensible data acquisition tool for Android'. In: (2016).
- [23] *Guide to app architecture*. Android Developer. URL: <https://developer.android.com/jetpack/docs/guide> (visited on 20/06/2019).
- [24] *Intents and Intent Filters*. Android Developer. URL: <https://developer.android.com/guide/components/intents-filters> (visited on 24/06/2019).
- [25] *Keep the device awake*. Android Developer. URL: <https://developer.android.com/training/scheduling/wakelock> (visited on 25/06/2019).

- [26] Walter L. Hursch and Cristina Videira Lopes. ‘Separation of Concerns’. In: (Mar. 1995), pp. 3, 16. URL: <http://citeseerx.ist.psu.edu/viewdoc/download;jsessionid=031E83D200FD8770C255B48EA0C2E1C2?doi=10.1.1.29.5223&rep=rep1&type=pdf> (visited on 24/05/2019).
- [27] Huoran Li, Xuanzhe Liu and Qiaozhu Mei. ‘Predicting Smartphone Battery Life based on Comprehensive and Real-time Usage Data’. In: (Jan. 2018), p. 2. URL: https://www.researchgate.net/publication/322498404_Predicting_Smartphone_Battery_Life_based_on_Comprehensive_and_Real-time_Usage_Data (visited on 30/05/2019).
- [28] *LiveData Overview*. Android Developer. URL: <https://developer.android.com/topic/libraries/architecture/livedata> (visited on 20/06/2019).
- [29] Tian Lou. ‘A Comparison of Android Native App Architecture – MVC, MVP and MVVM’. In: (), p. 57. URL: https://pure.tue.nl/ws/portalfiles/portal/48628529/Lou_2016.pdf (visited on 20/06/2019).
- [30] Innocent Mapanga and Prudence Kadebu. ‘Database Management Systems: A NoSQL Analysis’. In: *International Journal of Modern Communication Technologies And Research (IJMCTR)* Volume-1 (Sept. 2013), pp. 12–18. URL: https://www.researchgate.net/publication/258328266_Database_Management_Systems_A_NoSQL_Analysis (visited on 15/05/2019).
- [31] Stephen McGrath and Jonathan Whitty. ‘Stakeholder defined’. In: *International Journal of Managing Projects in Business* 10 (July 2017), pp. 4, 13, 14. DOI: 10.1108/IJMPB-12-2016-0097. (Visited on 01/05/2019).
- [32] Rajalakshmi Nandakumar, Shyamnath Gollakota and Nathaniel Watson. ‘Contactless Sleep Apnea Detection on Smartphones’. In: *GetMobile: Mobile Computing and Communications* 19 (Dec. 2015), pp. 22–24. DOI: 10.1145/2867070.2867078. (Visited on 26/06/2019).
- [33] T Penzel et al. ‘The use of a mobile sleep laboratory in diagnosing sleep-related breathing disorders’. In: *Journal of medical engineering & technology* 13.1-2 (1989), pp. 100–103. URL: <https://www.tandfonline.com/doi/pdf/10.3109/03091908909030206> (visited on 26/06/2019).
- [34] *Platform Architecture*. Android Developer. URL: <https://developer.android.com/guide/platform> (visited on 21/06/2019).
- [35] *Power Management*. Android Developer. URL: <https://source.android.com/devices/tech/power/mgmt> (visited on 25/06/2019).
- [36] *Processes and Threads*. Android Developer. URL: <https://stuff.mit.edu/afs/sipb/project/android/docs/guide/components/processes-and-threads.html#IPC> (visited on 24/06/2019).
- [37] *Save data in a local database using Room*. Android Developer. URL: <https://developer.android.com/training/data-storage/room/index> (visited on 20/06/2019).

- [38] *Services*. Android Developer. URL: <https://stuff.mit.edu/afs/sipb/project/android/docs/guide/components/services.html> (visited on 23/06/2019).
- [39] SweetZpot. *SweetZpot FlowTM Sensor*. URL: <https://www.sweetzpot.com/flow> (visited on 28/05/2018). URL: <https://www.sweetzpot.com/flow>.
- [40] *The color system*. Material Design. URL: <https://material.io/design/color/the-color-system.html> (visited on 04/06/2019).
- [41] Stein Kristiansen Thomas Peter Plagemann Vera Hermine Goebel. *Android based eHealth applications with BiTalino sensors*. URL: <http://www.mn.uio.no/ifi/studier-based-ehealth-applications-with-bitalino-s.html> (visited on 05/05/2018). 2015. URL: <http://www.mn.uio.no/ifi/studier/masteroppgaver/dmms/android-based-ehealth-applications-with-bitalino-s.html>.
- [42] M. U.Farooq et al. 'A Critical Analysis on the Security Concerns of Internet of Things (IoT)'. In: *International Journal of Computer Applications* 111.7 (18th Feb. 2015), pp. 1–6. ISSN: 09758887. DOI: 10.5120/19547-1280. URL: <http://research.ijcaonline.org/volume111/number7/pxc3901280.pdf> (visited on 30/05/2019).
- [43] Saurabh Zunke and Veronica D'Souza. 'JSON vs XML: A Comparative Performance Analysis of Data Exchange Formats'. In: 3.4 (2014), pp. 2–4. URL: <http://ijcsn.org/IJCSN-2014/3-4%20JSON-vs-XML-A-Comparative-Performance-Analysis-of-Data-Exchange-Formats.pdf> (visited on 30/05/2019).